

# Distributed Commit

# Distributed Commit Problem

---

- ◆ Some applications perform operations on multiple databases
  - For example, replicated DBs for fault tolerance
- ◆ We would like a guarantee that either all databases get updated, or none do
- ◆ **Distributed commit** problem
  - Operation is committed when all participants can perform it
  - Once a commit decision is reached, this requirement holds even if some participants fail and later recover

# ACID Properties

---

Transaction (group of operations) behaves as one operation

- ◆ **(Failure) Atomicity**: if transaction failed, then no changes apply to the database
- ◆ **Consistency**: database integrity constraints always hold
- ◆ **Isolation (Atomicity)**: partial results are hidden
- ◆ **Durability**: the effects of transactions that were committed are permanent

# 2PC Overview

---

- ◆ Assumes a **coordinator**
  - Initiates commit/abort
- ◆ Each database informs coordinator if it is ready to commit
  - Until the commit actually occurs, the update is considered temporary
  - A pending update can be discarded (ie, transaction aborts) until all databases say “ok”
- ◆ Coordinator decides outcome and informs all databases

**SOUNDS EASY!**

# 2PC (Simplified), No Failures

Coordinator:

Multicast *ready\_to\_commit*

Collect replies

All *Ok* => send *commit*


Else => send *abort*

Participant receives:

*ready\_to\_commit* => save to temp area and reply *Ok*


*commit* => make changes permanent

*abort* => delete temp area



# Participant States

- ◆ **Initial state:** ends when  $p_i$  received *ready\_to\_commit* and is ready to send its *Ok*
- ◆ **Prepared to commit:**  $p_i$  sent its *Ok*, saves in temp area and waits for the final decision (*commit* or *abort*) from coordinator
- ◆ **Commit or abort:**  $p_i$  knows the final decision, must execute it



# Participant Failures

---

- ◆ **Initial state:** if  $p_i$  crashes before receiving *ready\_to\_commit*, it does not send its *Ok* back, the coordinator will abort the protocol (not enough *Oks* are received)
- ◆ **Prepared to commit:** if  $p_i$  crashes before it learns the outcome, resources remained blocked. **It is critical that a crashed participant learn the outcome of pending operations when it comes back** (need logging system).
- ◆ **Commit or abort:** if  $p_i$  crashes before executing, it must complete the commit or abort repeatedly despite being interrupted by failures

# Dealing with Participant Failures

---

- ◆ A participant that crashed and recovered:
  - Must remember in what state it was before crashing
  - Must find out the (global) outcome by contacting the coordinator
- ◆ The coordinator:
  - Must keep track of pending protocol runs
  - Must find out when a participant indeed completed the decision



# Details

---

- ◆ Coordinator assigns unique identifiers for each protocol run
  - For example, use logical clocks + participant ID
- ◆ Messages carry the identifier of protocol run they are part of
- ◆ Lots of messages must be stored, need garbage collection, the challenge is to determine when it is safe to remove the information

# 2PC: Overcoming Participant Failures

---

Coordinator:

Multicast *ready\_to\_commit*

Collect replies

All *Ok* => log 'commit' to 'outcomes' table and send *commit*

Else => send *abort*

Collect acknowledgments

Garbage-collect protocol 'outcomes' information

Participant:

Receives:

*ready\_to\_commit* => save to temp area and reply *Ok*

*commit* => make changes permanent, send *acknowledgment*

*abort* => delete temp area

After recovering from failure:

For each pending protocol run: contact coordinator to learn outcome

# Coordinator Failures

---

- ◆ If coordinator crashes when collecting *Oks*:
  - Some participants will be ready to commit (sent *Ok*)
  - Others decided to abort, did not send *Ok*
  - Others may not know the state
- ◆ If coordinator crashes during its decision or before sending it out:
  - Some participants will be in prepare-to-commit state
  - Others will not know the outcome

# Modifications

---

- ◆ If coordinator fails, participants are blocked waiting for it to recover
- ◆ After the coordinator recovers, there are pending protocol runs that must be finished
- ◆ Coordinator must:
  - Remember its state before crashing
    - Write “commit” or “abort” into permanent storage before sending the decision to the other participants
  - Push pending operations through
- ◆ Participants may see duplicated messages

# 2PC Coordinator

---

Multicast *ready\_to\_commit*

Collect replies

All *OK* => log 'commit' to 'outcomes' table, wait until safely in persistent storage, then send commit

Else => send *abort*

Collect *acknowledgments*

Garbage collect protocol outcome information

After failure:

For each pending protocol run in 'outcomes' table

Send outcome (*commit* or *abort*)

Wait for *acknowledgments*

Garbage collect outcome information

# 2PC Participant

---

First time message received

*ready\_to\_commit*

save to temp area and reply *OK*

*commit*

make changes permanent

*abort*

delete temp area

Message is a duplicate (because of a recovering coordinator)

Send *acknowledgment*

After failure:

For each pending protocol run:

contact coordinator to learn outcome

# What If Coordinator Never Recovers?

---

- ◆ Instead of blocking, can allow remaining participants to complete the protocol on their own
  - Caveat: any participant taking over will not be able to safely conclude that the coordinator actually failed
- ◆ If timeout expires at a participant that is in the prepare-to-commit state:
  - The participant can send out the first-phase message, querying the other participants to learn outcome
  - Continue with second phase

# Allowing Progress

---

- ◆ Can a participant always determine the outcome?
  - Example: all participants are in prepared-to-commit state except  $p_j$ , which can not be reached
  - Only the coordinator and  $p_j$  can determine the outcome
- ◆ If the coordinator is itself a participant, only one failure blocks the protocol
- ◆ All participants must now maintain information about the outcome of the protocol until they are sure that all participants learned the outcome



# Garbage Collection

---

- ◆ Add another phase from the coordinator to all participants, telling participants that it is safe to garbage-collect the protocol information
- ◆ If coordinator fails:
  - If a participant is in the final state but did not see the garbage collect message, it will send again the commit or abort message
  - All participants acknowledge after they execute
  - Once all participants acknowledged the message, garbage collection message can be sent out
- ◆ Garbage collection can be run periodically

# 2PC: Coordinator (Final)

---

Multicast: ready\_to\_commit

Collect replies

All OK => log 'commit' to 'outcomes' table, wait until safe in persistent storage, send commit

Else => send abort

Collect acknowledgments

After failure:

For each pending protocol run in 'outcomes' table

Send outcome (commit or abort)

Wait for acknowledgments

Periodically

Query each participant about terminated protocol runs

Determine fully terminated protocol runs to garbage collect their outcome information

# 2PC: Participant (Final)

---

First time message received

*ready\_to\_commit*

save to temp area and reply *OK*

*commit*

Log outcome, make changes permanent

*abort*

Log outcome, delete temp area

Message is a duplicate (recovering coordinator)

Send *acknowledgment*

After failure:

For each pending protocol run:

contact coordinator to learn outcome

After timeout in prepare to commit state:

Query other participants about state

If outcome can be deduced: Run coordinator-recovery protocol

If outcome uncertain: must wait

# 2PC Summary

---

- ◆ Message complexity  $O(n^2)$
- ◆ Worst case: network disrupts the communication in each phase
- ◆ Pure 2PC will always block if coordinator fails
- ◆ Final version provides increased availability but can still block if a failure occurs at a critical stage: both coordinator and a participant fail during the decision stage
  - Trades liveness for safety

# 3-Phase Commit

---

- ◆ Guarantees that the protocol will not block when only fail-stop failures occur
  - Assumption: a participant fails only by crashing, crashes are accurately detectable
- ◆ Fundamental problem: coordinator makes a decision which will be known and acted upon by some participant, while other participants will not know it

# 3PC Key Idea

---

Additional round of communication (and delays) to ensure that the state of the system can always be deduced by a subset of alive participants that can communicate with each other

before the commit, coordinator tells all participants that everyone sent OKs

# 3PC, No Failures

Coordinator:

Multicast *ready\_to\_commit*

Collect OKs

All *Ok* => send *precommit*

Else => send *abort*

Collect ACKs

All *ACK* => send *commit*

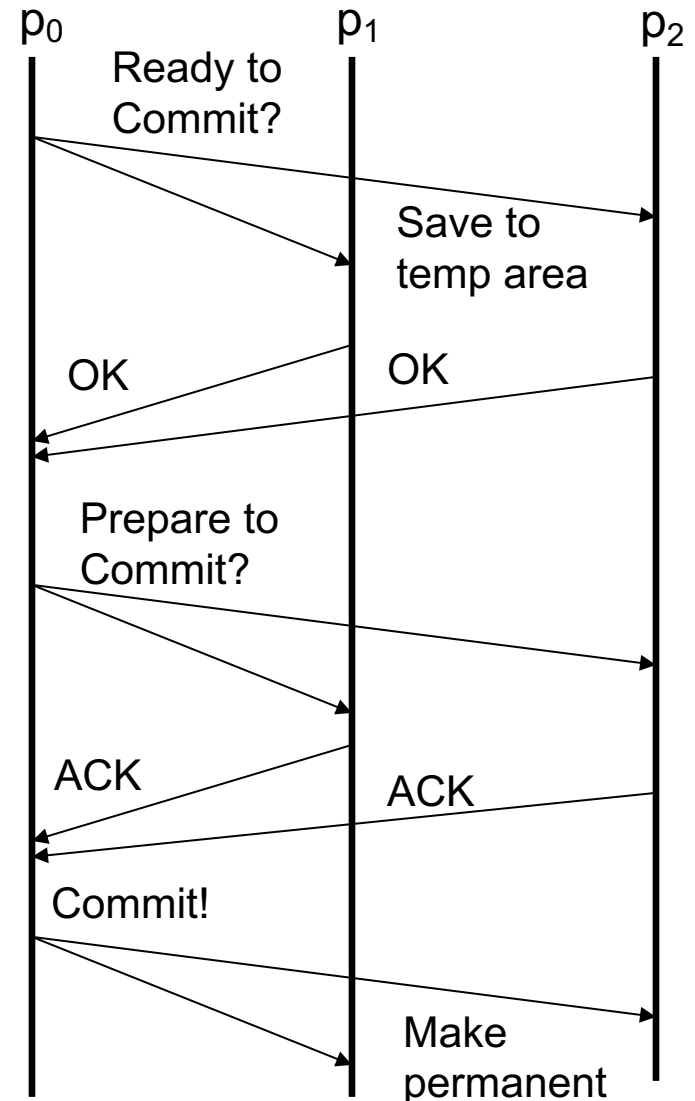
Participant receives:

*ready\_to\_commit* => save to temp area and reply *Ok*

*precommit* => send *ACK*

*commit* => make changes permanent

*abort* => delete temp area



# Dealing with Failures

---

- ◆ Alive participants select a new coordinator and try to complete transaction, based on their current states
  - Membership is static, detection is accurate, alive participant with lowest id is selected
- ◆ If crashed nodes committed or aborted, then survivors should not contradict; otherwise, survivors can do as they decide



# 3PC: Coordinator

---

Multicast *ready\_to\_commit*

Collect replies

All *OK* => log 'precommit' and send precommit

Else => send *abort*

Collect acks from alive participants

All *ack* => log commit and send commit

Collect acknowledgements that operation was finished

Garbage-collect protocol outcome information

# 3PC: Participant

---

Participant logs its state after each message

*ready\_to\_commit*

save to temp area and reply *OK*

*precommit*

Enter precommit state, send *ack*

*commit*

make changes permanent

*abort*

delete temp area

After failure:

Collect participant state information

All *precommit* or any *commit*, push forward the commit

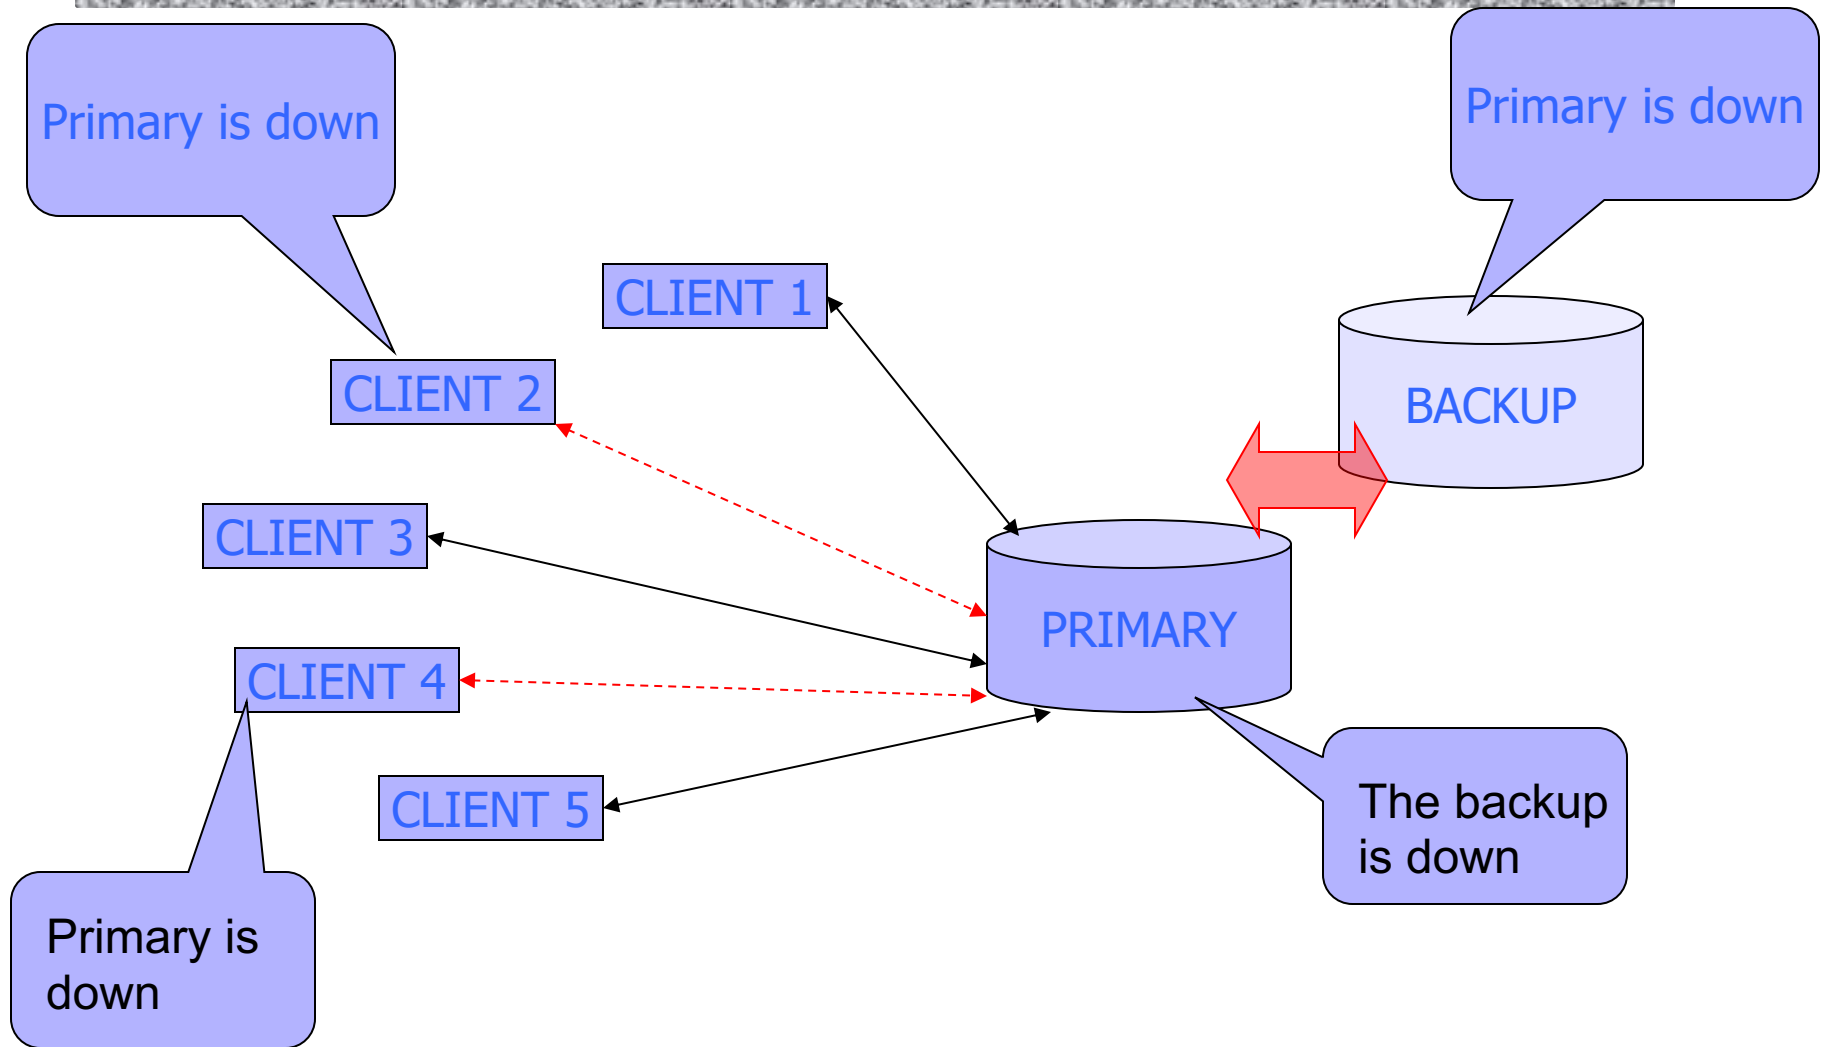
Else, push back the abort

# 3PC and Network Partitions

---

- ◆ Suppose a network partition separates the participants in two groups:
  - One group has received precommit, so they all terminate the protocol by commit
  - The other group sees a state that is ok to commit but they have not received precommit, so they all abort upon timeout
- ◆ 3PC does not work in case of network partitions

# Things Goes Wrong...



# 3PC Summary

---

- ◆ Requires 3 phases
  - 4 with garbage collection
- ◆ Works only under fail-stop failures
- ◆ Does not work (produces an inconsistent global state) if network partitions happen
  - 3PC trades safety for liveness
  - Remember CAP theorem?
- ◆ Can we design a protocol that is **safe and live**?
  - FLP impossibility result