


Time Global State Failure Detection


Vitaly Shmatikov

Time

- ◆ Time is essential for ordering events in a distributed system
 - **Physical time:** local clock, global clock
 - **Logical time:** Lamport clocks, vector clocks




Historical Clocks




Electrical Clocks

- ◆ First developed in 1920s
 - Uses carefully shaped quartz crystal
 - Pass current, counts oscillations
- ◆ Most oscillate at 32,768/sec
 - Easy to count in hardware
 - Small enough to fit ($\sim 4\text{mm}$)
- ◆ Typical quartz clock quite accurate
 - Within 15 sec/30 days ($6\text{e-}6$)
 - Can achieve $1\text{e-}7$ accuracy in controlled conditions
 - Not good enough for today's applications



Atomic Clocks

- ◆ Based on atomic physics
 - Cool atoms to near absolute zero
 - Bombard them with microwaves
 - Count transitions between energy levels
- ◆ Most accurate timekeeping devices
 - Accurate to within 10^{-9} seconds per day (loses 1 second in 30 million years)
- ◆ Standard International Second defined in terms of atomic oscillations
 - 9,192,631,770 transitions of cesium-133 atom



International Atomic Time

- ◆ Atomic clocks used to define several time standards
- ◆ TAI: International Atomic Time
 - Avg. of 200 atomic clocks, corrected for time dilation
 - Essentially, a count of the number of seconds passed since January 1, 1958
- ◆ UTC: since January 1, 1972, defined to follow TAI with an exact offset of an integer number of seconds, changing only when a leap second is added to keep clock time synchronized with the rotation of the Earth

Using Real Clocks to Order Events

- ◆ Each event carries a timestamp
- ◆ **Global clock:** processes have access to a central global clock
 - The global clock gives global ordering of events
- ◆ **Local clock:** each process has its own clock
 - What if the clocks are not synchronized?
 - What if events happened at the same time?

Clocks in Computers

- ◆ Real-time clock: CMOS clock (counter) circuit driven by a quartz oscillator with battery backup to continue measuring time when power is off
- ◆ OS generally programs a timer circuit to generate an interrupt periodically
 - e.g., 60, 100, 250, 1000 interrupts per second
 - Programmable Interval Timer (PIT) – Intel 8253, 8254
 - Interrupt handler adds 1 to a counter in memory
- ◆ Quartz oscillators oscillate at slightly different frequencies, clocks do not agree in general

When Is a Clock “Correct”?

◆ Relative to an “ideal” clock

- Clock skew is magnitude
- Clock drift is difference in rates

◆ Say clock is correct within p if

$$(1-p)(t'-t) \leq H(t') - H(t) \leq (1+p)(t'-t)$$

- $(t'-t)$ True length of interval
- $H(t') - H(t)$ Measured length of interval
- $(1-p)(t'-t)$ Smallest acceptable measurement
- $(1+p)(t'-t)$ Largest acceptable measurement

◆ Monotonic property: $t < t' \Rightarrow H(t) < H(t')$

Monotonicity

- ◆ If a clock is running “slow” relative to real time...
 - Can simply re-set the clock to real time
 - Doesn't break monotonicity
- ◆ What if a clock is running “fast”?
 - Re-setting the clock back breaks monotonicity
 - Imagine programming with the same time occurring twice
- ◆ Instead, “slow down” clock
 - Maintains monotonicity

Network Time Protocol (NTP)

◆ NTP is a distributed service that...

- Keeps machines synchronized to UTC
- Deals with lengthy losses of connectivity
- Enables clients to synchronize frequently (scalable)
- Avoids security attacks

◆ NTP deployed widely today

- Uses 64-bit value, epoch is 1/1/1900 (rollover in 2036)
- Precision: 1ms on LANs, 10s of ms on Internet
- **NTP pool** is a dynamic collection of 4000 servers that volunteer to provide time via NTP


Reference Clocks

- ◆ Many NTP servers synchronize directly to UTC using specialized equipment
 - **Atomic clocks:** ultimately are the root source of time in NTP
 - **Global Positioning System (GPS):** can synchronize with a satellite's atomic clock
 - **Code Division Multiple Access (CDMA):** can synchronize with a local wireless provider (who in turn most likely synchronizes using GPS)
 - **Radio signals:** similar to CDMA, can synchronize with time/frequency radio stations

From Physical to Logical Clocks

- ◆ Synchronized clocks are great if we have them
- ◆ Why do we need the time anyway?
- ◆ In distributed systems, we care about “what happened before what”
- ◆ Message-based systems, two type of events
 - Send a message
 - Receive a message

"Happened Before"




- ◆ If events a and b take place at the same process and a occurs before b (physical time), then we have $a \rightarrow b$
- ◆ If a is a send event of message m at p_1 and b is a deliver event of the same m at p_2 , $p_1 \neq p_2$ then $a \rightarrow b$
- ◆ If $a \rightarrow b$ and $b \rightarrow c$ then $a \rightarrow c$

Reminder: Partial and Total Order


- ◆ A relation R over a set S is a **partial order** iff for each a, b , and c in S :
 - aRa (reflexive)
 - $aRb \wedge bRa \Rightarrow a = b$ (antisymmetric)
 - $aRb \wedge bRc \Rightarrow aRc$ (transitive)
- ◆ A relation R over a set S is **total order** if for each distinct a and b in S , R is antisymmetric, transitive and either aRb or bRa (completeness)

Lamport Clocks (1978)



- ◆ Each process maintains its own clock C_i (a counter)
- ◆ For any events a and b in process p_i
if $a \rightarrow b$ then $C_i(a) < C_i(b)$
- ◆ Implementation:
 - Each p_i increments C_i between any successive events
 - On sending a message m , attach local clock $T_m = C_i(a)$
 - On receiving a message m , process p_k sets C_k to $C_k = \max(C_k, T_m) + 1$

Lamport Clocks: Example



Lamport Clocks: Total Order

- ◆ Logical clocks only provide partial order
- ◆ Create total order by breaking the ties
- ◆ Example: have an order on process identifiers, use them to break ties
 - If a is event in p_i and b is event in p_j then
$$a \rightarrow b \quad \text{iff}$$
 - $C_i(a) < C_j(b)$ or
 - $C_i(a) = C_j(b)$ and $p_i < p_j$

Concurrent Events


- ◆ If $a \rightarrow b$ and $b \rightarrow a$ then a and b are **concurrent**
- ◆ Logical clocks assign order to events that are causally independent
 - Events that are causally independent appear as if they happened in a certain order
- ◆ For some applications (e.g. debugging) it is important to capture independence

Vector Clocks

 [Fidge and Mattern (independently), 1988]

- ◆ Each process p_i maintains a vector C_i
- ◆ When p_i executes an event, it increments its own clock $C_i[i]$
- ◆ When p_i sends a message m to p_j , it attaches its vector C_i
- ◆ When p_i receives a message m , increments its own clock and updates the clock for the other processes
$$\forall j: 1 \leq j \leq n, j \neq i: C_i[j] = \max(C_i[j], m.C[j])$$
$$C_i[i] = C_i[i] + 1.$$

Vector Clocks: Example



How to Order with Vector Clocks


Given two events a and b , $a \rightarrow b$ if and only if $V(a)$ is less than or equal to $V(b)$ for all process indices, and at least one of those relationships is strictly smaller

- $a \rightarrow b \equiv \forall i: 1 \leq i \leq n: V(a)[i] \leq V(b)[i] \wedge \exists i: 1 \leq i \leq n: V(a)[i] < V(b)[i]$

Otherwise, they are concurrent or independent

- $a \parallel b \equiv \exists i: 1 \leq i \leq n: V(a)[i] < V(b)[i] \wedge \exists j: 1 \leq j \leq n: V(b)[j] < V(a)[j]$

What Events Are Independent?



Why We Need Global Snapshots

- ◆ Checkpointing: save the state and restart the distributed application after a failure
- ◆ Garbage collection of objects: objects at servers that don't have any other objects (at any servers) with pointers to them
- ◆ Deadlock detection: debugging for database transaction systems
- ◆ Termination of computation: useful for batch computing systems

Recording Global Snapshots

- ◆ If synchronized clocks are available, each process records its state at a known time t
 - How to obtain the state of the messages that transit the channels?
- ◆ If synchronized clocks are not available?
 - How to determine when a process takes its snapshot?
 - How to distinguish between the messages to be recorded in the snapshot from those not to be recorded?

Chandy-Lamport Algorithm

records a consistent global state of an asynchronous system.

◆ System model

- No failures and all messages arrive intact and only once
- Communication channels are unidirectional and FIFO
- There is a communication path between any two processes

◆ Other assumptions

- Any process may initiate the snapshot algorithm
- The snapshot algorithm does not interfere with the normal execution of the processes
- Each process records its local state and the state of its incoming channels

Chandy-Lamport Algorithm

◆ A process needs to know

- When to start recording (if it was not the one that initiated the algorithm)
- What messages to include in the snapshot
- When did all other processes record their snapshot

◆ Key design: a control message, **marker**

- To separate messages to be included from messages not to be included
- To inform other processes that it has recorded snapshot
- To tell other processes to start recording: a process must record its snapshot no later than when it receives a marker on any of its incoming channels

Chandy-Lamport Algorithm

- ◆ Any process can initiate by executing the “Marker Sending Rule”
- ◆ On receiving a marker on some channel c , a process executes the “Marker Receiving Rule”
 - If not yet recorded its local state, record the state of channel c as empty and execute the “Marker Sending Rule” to record its local state
- ◆ The algorithm terminates after each process has received a marker on all of its incoming channels
- ◆ Sum of all local snapshots = global state

Chandy-Lamport Snapshot Algorithm

- ◆ Marker-sending rule for a process
 - Record its local state
 - Send a marker to all other processes on the corresponding channels before sending any other message
- ◆ Marker-receiving rule for a process q on channel c
 - If not yet recorded its local state, then
 - Record its local state
 - Record the state of channel c as “empty”
 - Turn on recording of messages over other incoming channels
 - Send a marker on each outgoing channel
 - Else
 - Record the state of incoming channel as all messages received over it after q recorded its state and before it received the marker along c

Chandy-Lamport Algorithm: Example


◆ Setup

- Three processes p, q and r
- Communication channels: c1 (p to q), c2 (q to p), c3 (q to r), and c4 (r to p)
- All start with state = \$500 and the channels are empty. The stable property is that the total amount of money is \$1500.

◆ Process p sends \$10 to q and then starts the snapshot algorithm: records its current state \$490 and sends out a marker on c1


◆ Meanwhile q has sent \$20 to p along c2 and 10 to r along c3

Snapshot/State Recording Example (Step 1)




Node	Recorded State				
	state	c1	c2	c3	c4
p	490		{}		{}
q		{}			
r				{}	

Snapshot/State Recording Example (Step 2)




Node	Recorded State				
	state	c1	c2	c3	c4
p	490		{}		{}
q	480	{}			
r				{}	

Snapshot/State Recording Example (Step 3)




Node	Recorded State				
	state	c1	c2	c3	c4
p	490		{}		{}
q	480	{}			
r	485			{}	

Snapshot/State Recording Example (Step 4)



Node	Recorded State				
	state	c1	c2	c3	c4
p	490		{20}		{}
q	480	{}			
r	485			{}	

Snapshot/State Recording Example (Step 5)



Node	Recorded State				
	state	c1	c2	c3	c4
p	490		{20}		{25}
q	480	{}			
r	485			{}	

Correctness for Chandi-Lamport

- ◆ How do we define correctness?
- ◆ Records a **consistent global state** of an asynchronous system

History of Events

Given a process p_i

◆ e_i^j is the j^{th} event

◆ History is a sequence of events

$$h_i = \langle e_i^0, e_i^1, \dots \rangle$$

◆ Prefix history is the history up to the k^{th} event

$$h_i^k = \langle e_i^0, e_i^1, \dots, e_i^k \rangle$$

◆ State S_i^k is the state of process p_i immediately before the k^{th} event

More Definitions

- ◆ Global history: the set of all processes' histories

$$H = \cup_i (h_i)$$

- ◆ Global state: the set of all processes' states

$$S = \cup_i (S_i^{k_i})$$

- ◆ Cut: a set of prefix histories

$$C \subseteq H = h_1^{c1} \cup h_2^{c2} \cup \dots \cup h_n^{cn}$$


- ◆ Frontier of a cut: the set of last events that happened in each prefix history

$$C = \{e_i^{ci}, i = 1, 2, \dots n\}$$

Consistent Cuts

A cut C is consistent if for any event e in the cut, if an event f 'happened before' e , then f is also in the cut C

$$\forall e \in C (\text{if } f \rightarrow e \text{ then } f \in C)$$



Using Global States

- ◆ **Consistent global state:** a global state that corresponds to a consistent cut
- ◆ **Run:** a total ordering of events in history H that is consistent with each process history h_i
- ◆ **Linearization:** a run consistent with happens-before relation in H ; **linearizations pass through consistent global states**
- ◆ **Reachability:** a global state S_k is reachable from global state S_i , if there is a linearization, L , that passes through S_i and then through S_k

Global State Predicates

- ◆ Global state predicate: a function from the set of global states to $\{\text{TRUE}, \text{FALSE}\}$
- ◆ Stable global state predicate: one that once it becomes true, it remains true in all future states reachable from that state
- ◆ Examples:
 - “the system is deadlocked”
 - “all tokens in a token ring have disappeared”
 - “the computation has finished”

Safety and Liveness

- ◆ **Safety:** a condition that must hold in every finite prefix of a sequence (from an execution)
“nothing bad happens”
- ◆ **Liveness:** a condition that must hold a certain number of times
“something good happens”

Stable Global States and Safety

Assume that a “bad thing” BT (for example deadlock) is a global state predicate and S_0 is the initial state of the system, then

“Safety with respect to BT” means

$\forall S$ reachable from S_0 , $BT(S) = \text{FALSE}$

Stable Global States and Liveness

Assume that a “good thing” GT (for example, reaching termination) is a global state predicate and S_0 is the initial state of the system, then

Liveness with respect to GT means

For any linearization L starting at S_0 \exists state S_L reachable from S_0 such that $GT(S_L) = \text{TRUE}$

Failure Detectors as Abstraction

- ◆ **Failure detector:** distributed oracle that makes guesses about process failures
 - Failure = crash (only for now!)
- ◆ Accuracy: failure detector makes no mistakes when labeling processes as crashed
- ◆ Completeness: failure detector “eventually” (after some time) suspects every process that actually crashed
- ◆ Used to solve different distributed systems problems

Completeness

- ◆ Strong completeness: There is a time after which every process that crashes is suspected by EVERY correct process
- ◆ Weak completeness: There is a time after which every process that crashes is suspected by SOME correct process

Accuracy

- ◆ Strong accuracy: No process is suspected before it crashes
- ◆ Weak accuracy: At least one correct process is never suspected
- ◆ Eventual strong accuracy: There is a time after which correct processes are not suspected by any correct process
- ◆ Eventual weak accuracy: There is a time after which some correct process is never suspected by any correct process

Perfect Failure Detector

- ◆ A perfect failure detector has strong accuracy and strong completeness... such detector is IMPOSSIBLE
- ◆ We have to live with unreliable failures detectors



Failure Detection Implementation

- ◆ **Push:** processes keep sending heartbeats “I am alive” to the monitor. If no message is received for awhile from some process, that process is suspected as being dead (faulty)
- ◆ **Pull:** monitor asks the processes “Are you alive?”, and process will respond “Yes”. If no answer is received from some process, the process is suspected as being dead (faulty)
- ◆ What are advantages and disadvantages of these two approaches?

Failure Detection Implementation

- ◆ Every process must know about who failed
- ◆ How to disseminate the information?
- ◆ What if not every node can communicate directly with another node?
 - Centralized
 - All-to-all
 - Gossip based: provides probabilistic guarantees

Metrics for failure detectors

- ◆ Detection time
- ◆ Mistake recurrence time
- ◆ Mistake duration
- ◆ Average mistake rate
- ◆ Query accuracy probability
- ◆ Good period duration
- ◆ Network load

