


TensorFlow

Goals

- ◆ Library for numerical computations (read: ML)
 - Primitives for defining functions on tensors and automatically computing their derivatives
 - **Tensor** \sim multi-dimensional array of numbers
- ◆ Specify computation as a data-flow graph
 - Nodes: operations with any number of tensor inputs and tensor outputs
 - Edges: tensors that flow between operations

Dataflow Programming

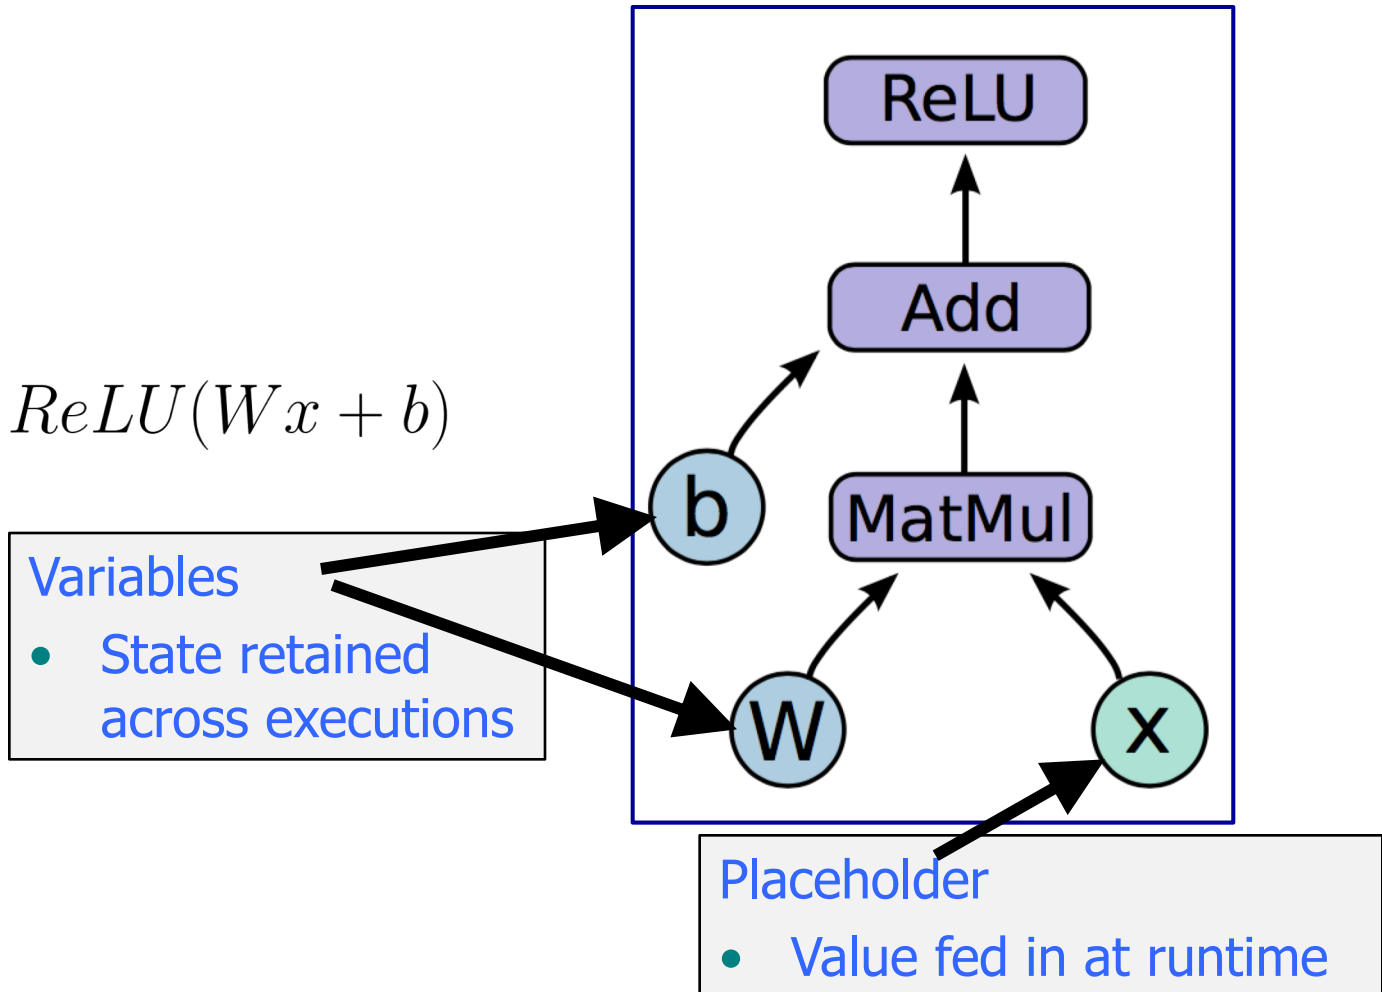


All nodes return tensors
TensorFlow does not care how
a node computes

```
define { node1 = tf.constant(3.0, dtype=tf.float32)
        node2 = tf.constant(4.0, dtype=tf.float32)
        node3 = tf.add(node1, node2)
exec { tf.Session().run(node3) #returns 7 }
```


Example

$$h = \text{ReLU}(Wx + b)$$




In Code

```
import tensorflow as tf
b = tf.Variable(tf.zeros((100,)))
W = tf.Variable(tf.random_uniform((784, 100), -1, 1))
x = tf.placeholder(tf.float32, (100, 784))
h = tf.nn.relu(tf.matmul(x, W) + b)
```



Run Graph in Execution Environment


```
import numpy as np
import tensorflow as tf
b = tf.Variable(tf.zeros((100,)))
W = tf.Variable(tf.random_uniform((784, 100), -1, 1))
x = tf.placeholder(tf.float32, (100, 784))
h = tf.nn.relu(tf.matmul(x, W) + b)
```



```
sess = tf.Session()
sess.run(tf.initialize_all_variables())
sess.run(h, {x: np.random.random(100, 784)})
```

Graph executes here

Partial Execution of Subgraphs




Parameter Server

- Hold mutable state
- Apply updates
- Maintain availability

server nodes:

worker nodes:



- Perform computation
- Mostly stateless (can checkpoint state to file system)
- Can be restarted

Parameter Server Example

```
with tf.device("/jobs:ps/task:0/cpu:0"):
    W = tf.Variable(...)
    b = tf.Variable(...)

inputs = tf.split(0,num_workers,input)
outputs = []

for i in range (num_workers):
    with tf.device("/job:worker/task:%d/gpu:0" % i):
        outputs.append(tf.matmul(input[i],W) + b)
```

Computing Gradients

TensorFlow nodes have attached gradient operations

```
tf.train.GradientDescentOptimizer(...).minimize(...)
```

adds optimization operation to computation graph

Gradients with respect to parameters are computed automatically via backpropagation

Training an ML Model

```
prediction = tf.nn.softmax(...)
label = tf.placeholder(tf.float32, [None, 10])

cross_entropy = tf.reduce_mean(-tf.reduce_sum(label *
tf.log(prediction), reduction_indices=[1]))

train_step =
tf.train.GradientDescentOptimizer(0.5).minimize(cross_entropy)

sess = tf.Session()
sess.run(tf.initialize_all_variables())


for i in range(1000):
    batch_x, batch_label = data.next_batch()
    sess.run(train_step, feed_dict={x: batch_x, label: batch_label})
```

Ex: Linear Regression in TensorFlow (1)

```
import numpy as np
import seaborn

# Define input data
X_data = np.arange(100, step=.1)
y_data = X_data + 20 * np.sin(X_data/10)

# Plot input data
plt.scatter(X_data, y_data)
```



Ex: Linear Regression in TensorFlow (2)

```
# Define data size and batch size
```

```
n_samples = 1000
```

```
batch_size = 100
```

```
# Tensorflow is finicky about shapes, so resize
```

```
X_data = np.reshape(X_data, (n_samples,1))
```

```
y_data = np.reshape(y_data, (n_samples,1))
```

```
# Define placeholders for input
```

```
X = tf.placeholder(tf.float32, shape=(batch_size, 1))
```

```
y = tf.placeholder(tf.float32, shape=(batch_size, 1))
```

Ex: Linear Regression in TensorFlow (3)

```
# Define variables to be learned
with tf.variable_scope("linear-regression"):
    W = tf.get_variable("weights", (1, 1),
                        initializer=tf.random_normal_initializer())
    b = tf.get_variable("bias", (1,),
                        initializer=tf.constant_initializer(0.0))
    y_pred = tf.matmul(X, W) + b
    loss = tf.reduce_sum((y - y_pred)**2/n_samples)
```

Note `reuse=False` so these tensors are created anew

$$J(W, b) = \frac{1}{N} \sum_{i=1}^N (y_i - (Wx_i + b))^2$$


Ex: Linear Regression in TensorFlow (4)

Sample code to run one step of gradient descent

```
In [136]: opt = tf.train.AdamOptimizer()
```

```
In [137]: opt_operation = opt.minimize(loss)
```

Note TensorFlow scope is not python scope! Python variable `loss` is still visible.



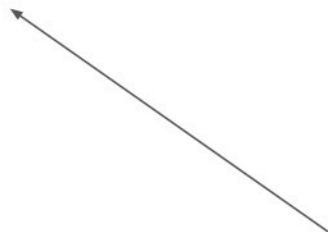
```
In [138]: with tf.Session() as sess:
```

```
.....:     sess.run(tf.initialize_all_variables())
```

```
.....:     sess.run([opt_operation], feed_dict={X: X_data, y: y_data})
```

```
.....:
```

But how does this actually work under the hood? Will return to TensorFlow computation graphs and explain.

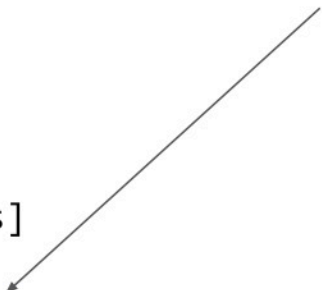


Ex: Linear Regression in TensorFlow (4)

```
# Sample code to run full gradient descent:  
# Define optimizer operation  
opt_operation = tf.train.AdamOptimizer().minimize(loss)
```

```
with tf.Session() as sess:  
    # Initialize Variables in graph  
    sess.run(tf.initialize_all_variables())  
    # Gradient descent loop for 500 steps  
    for _ in range(500):  
        # Select random minibatch  
        indices = np.random.choice(n_samples, batch_size)  
        X_batch, y_batch = X_data[indices], y_data[indices]  
        # Do gradient descent step  
        _, loss_val = sess.run([opt_operation, loss], feed_dict={X: X_batch, y: y_batch})
```

*Let's do a deeper.
graphical dive into
this operation*



Ex: Linear Regression in TensorFlow (5)

