

Spanner

Spanner

- ◆ Scalable, multi-version, globally distributed, synchronously replicated database
 - Externally consistent distributed transactions
 - Lock-free read-only transactions
 - Atomic schema changes
- ◆ Scale up to
 - millions of machines
 - hundreds of datacenters
 - trillions of database rows

Linearizability

A guarantee for a single operation on a single object

Informally...

- ◆ Writes should appear instantaneously within a system
- ◆ All later (by wall clock) reads reflect a value written at this or later time

Serializability (Isolation)

A guarantee for transactions consisting of one or more operations on one or more objects

- ◆ A set of transactions should execute as though each transaction ran in some serial order
- ◆ No deterministic order (no wall-clock constraints)
- ◆ This is isolation in ACID properties


Strict Serializability

- ◆ Linearizability + serializability
- ◆ Transactions have some serial behavior corresponding to wall-clock time
- ◆ Straightforward for non-overlapping transactions but what about **concurrent transactions**?

Concurrency Controls

- ◆ **Global lock:** simple, slow
- ◆ **Per-object lock:** doesn't guarantee serializability (isolation)

Partitioned Data Over Servers



- ◆ Why not just use 2-phase locking?
 - Grab locks over entire read and write set
 - Perform writes
 - Release locks (at commit time)

Concurrency Controls

- ◆ **Global lock:** simple, slow
- ◆ **Per-object lock:** doesn't guarantee serializability (isolation)
- ◆ **2-phase locking:** serializability, but can deadlock
- ◆ **Optimistic concurrency control:** good if few conflicts, bad performance if many conflicts
- ◆ **Multi-version concurrency control:** snapshot isolation (weaker than serializability)

DISRUPTIVE IDEA


Do clocks **really** need to be
arbitrarily unsynchronized?

Can you engineer some max divergence?

Key Idea Behind Spanner

- ◆ Attach global commit timestamps to transactions, even though transactions may be distributed
 - Timestamps represent serialization order: if transaction T1 commits before another transaction T2 starts, then T1's commit timestamp is smaller than T2's
- ◆ How to get the global timestamps: **TrueTime**
- ◆ Use existing algorithms such as Paxos and 2PC

Partitioned Data Over Servers



◆ How do you get serializability?

- Single machine: single COMMIT op in write-ahead log
- Distributed setting: assign global timestamp to txn (at some time after lock acquisition and before commit)
 - Centralized transaction manager
 - Distributed consensus on timestamps (not all ops)


Google's Setting

- ◆ Dozens of zones (datacenters)
- ◆ Per zone, 100-1000s of servers
- ◆ Per server, 100-1000 partitions (tablets)
- ◆ Every tablet replicated for fault-tolerance (e.g., 5x)

Spanner Features

- ◆ Applications can control replication configurations for data, specify constraints
 - which datacenters contain which data, how far data is from its users (to control read latency)
 - how far replicas are from each other (to control write latency)
 - how many replicas are maintained (to control durability, availability, and read performance)
- ◆ Data can also be dynamically and transparently moved between datacenters by the system to balance resource usage across datacenters

Architecture



◆ Universe master

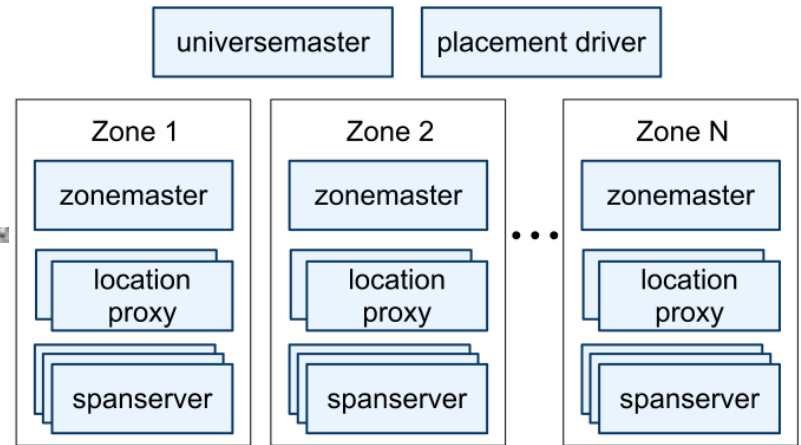
◆ Placement driver

- Handles automated movement of data across zones on the timescale of minutes
- Periodically communicates with the spanservers to find data that needs to be moved, either to meet updated replication constraints or to balance load.

◆ Universe consists of zones

- Denotes physical isolation
- Several zones can be in a datacenter

Zones



◆ Zonemaster

- Assigns the data to spanservers


◆ Spanservers

- Hundreds to thousands
- Store data
- Responsible for 100-1000 instances of a data structure called a **tablet** (different from the BigTable tablet)

◆ Location proxies

- Used by clients to locate the spanservers assigned to serve their data

Scale-out vs. Fault Tolerance



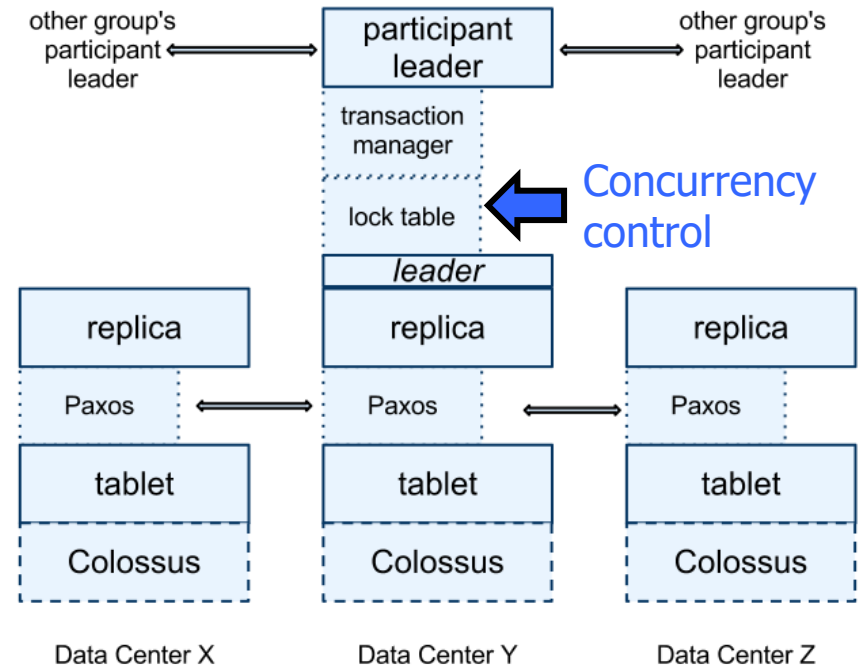
- ◆ Every tablet replicated via Paxos (with leader election)
- ◆ So every “operation” within transactions across tablets is actually a replicated operation within Paxos RSM
- ◆ Paxos groups can stretch across datacenters!

Tablet and Directory

- ◆ **Tablet** implements a bag of mappings
(key: string, timestamp: int64) → string
 - Spanner assigns **timestamps** to data... unlike Bigtable, Spanner is more like a multi-version database than a key-value store
- ◆ **Directory** is a set of contiguous keys that share a common prefix
 - Smallest unit of data placement
 - Smallest unit to define replication properties

Replication

- ◆ Each tablet replicated using Paxos
 - Stores metadata and logs of the tablet
- ◆ Leader among replicas in a Paxos group is chosen and all write requests for replicas in that group start at the leader



Paxos Leader Leases

- ◆ Paxos uses timed leases to make leadership long-lived (10 seconds by default)
 - Potential leader sends requests for timed lease votes
 - Upon receiving a quorum of lease votes the leader knows it has a lease
 - A replica extends its lease vote implicitly on a successful write, and the leader requests lease-vote extensions if they are near expiration
- ◆ Spanner depends on (and enforces) invariant:
for each Paxos group, each Paxos leader's lease interval is disjoint from every other leader's


Paxos Leaders and Time

- ◆ A Paxos leader can abdicate by releasing replicas from their lease votes. To preserve the disjointness invariant, Spanner constrains when abdication is permissible.
 - S_{\max} is the maximum timestamp used by a leader
 - Leader uses $\text{after}(S_{\max})$ to check if S_{\max} is passed so it can abdicate and release its replicas
- ◆ Paxos leaders cannot assign timestamp S_i greater than S_{\max} for transaction T_i and clients cannot see the data committed by transaction T_i till $\text{after}(S_i)$ is true
- ◆ Replicas maintain a timestamp t_{safe} which is the maximum timestamp at which that replica is up to date

All of this depends on global timestamps

TrueTime

- ◆ “Global wall-clock time” with bounded uncertainty
 - Timestamps become intervals, not single values



Consider event e_{now} which invoked $tt = \text{TT.now}()$:

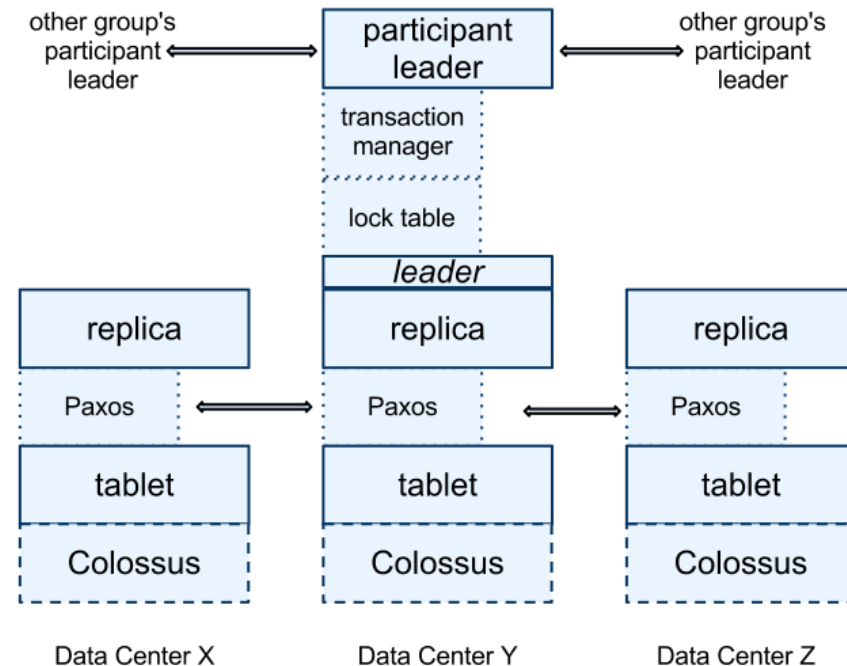
Guarantee: $tt.\text{earliest} \leq t_{\text{abs}}(e_{\text{now}}) \leq tt.\text{latest}$

TrueTime


- ◆ Daemon polls variety of timemasters and reaches consensus about correct timestamp
 - Majority of time masters are GPS-fitted, a few "Armageddon masters" are atomic clock-fitted
 - Different failure rates and scenarios
- ◆ TrueTime API
 - Key method is **now()** which not only returns current system time but also the **maximum uncertainty ϵ** (less than 10ms) in the time returned
 - **After(t)** – returns TRUE if t is definitely passed
 - **Before(t)** – returns TRUE if t is definitely not arrived

Transaction Manager

- ◆ If a transaction involves more than one Paxos group, those groups' leaders coordinate to perform two-phase commit
- ◆ One of the tablet groups is chosen as the coordinator
- ◆ The state of each transaction manager is stored in the underlying Paxos group (and therefore is replicated)

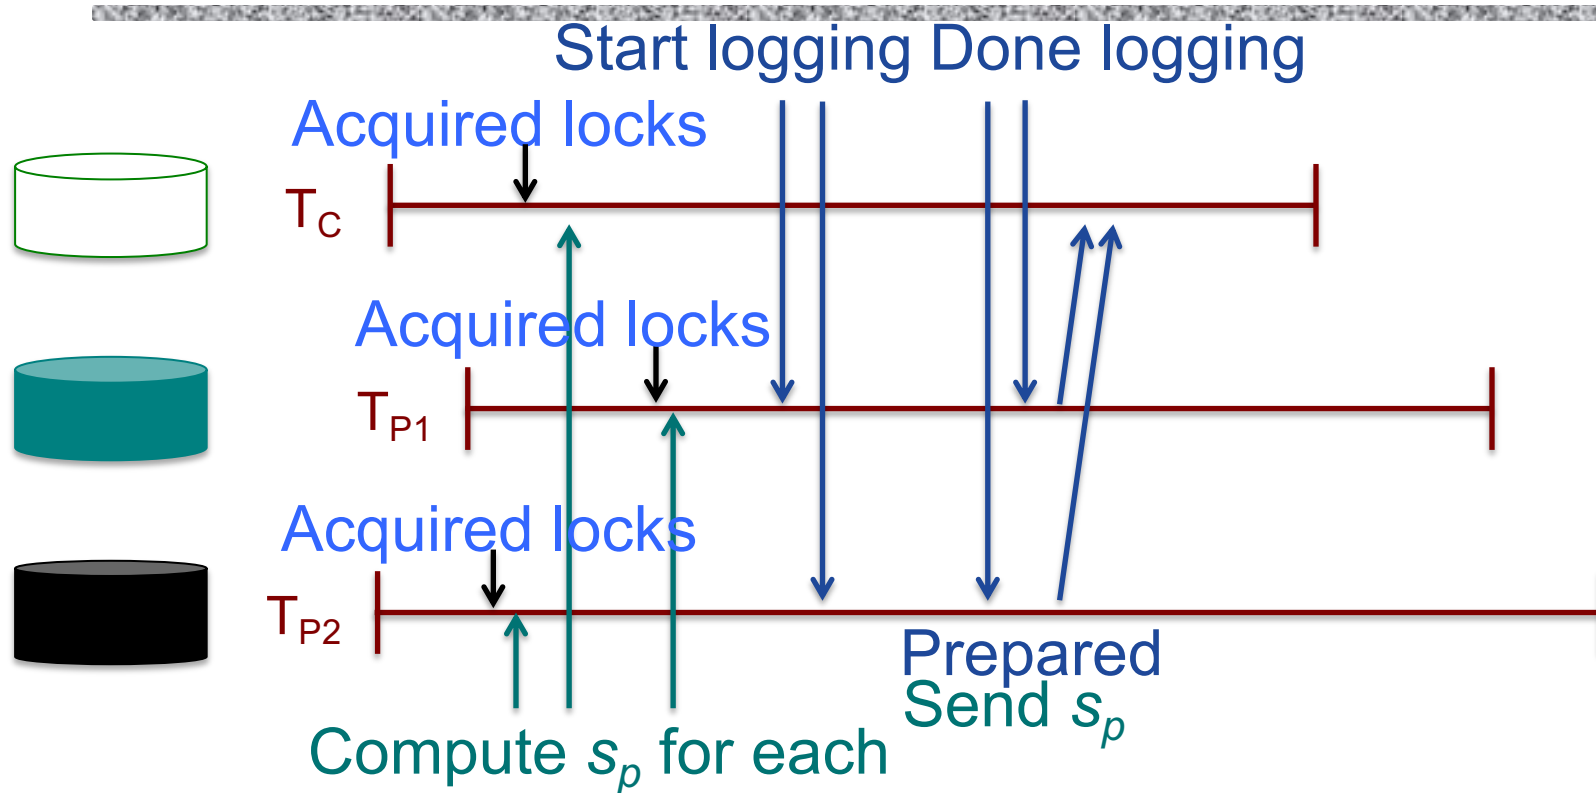


Commit Wait and 2-Phase Commit



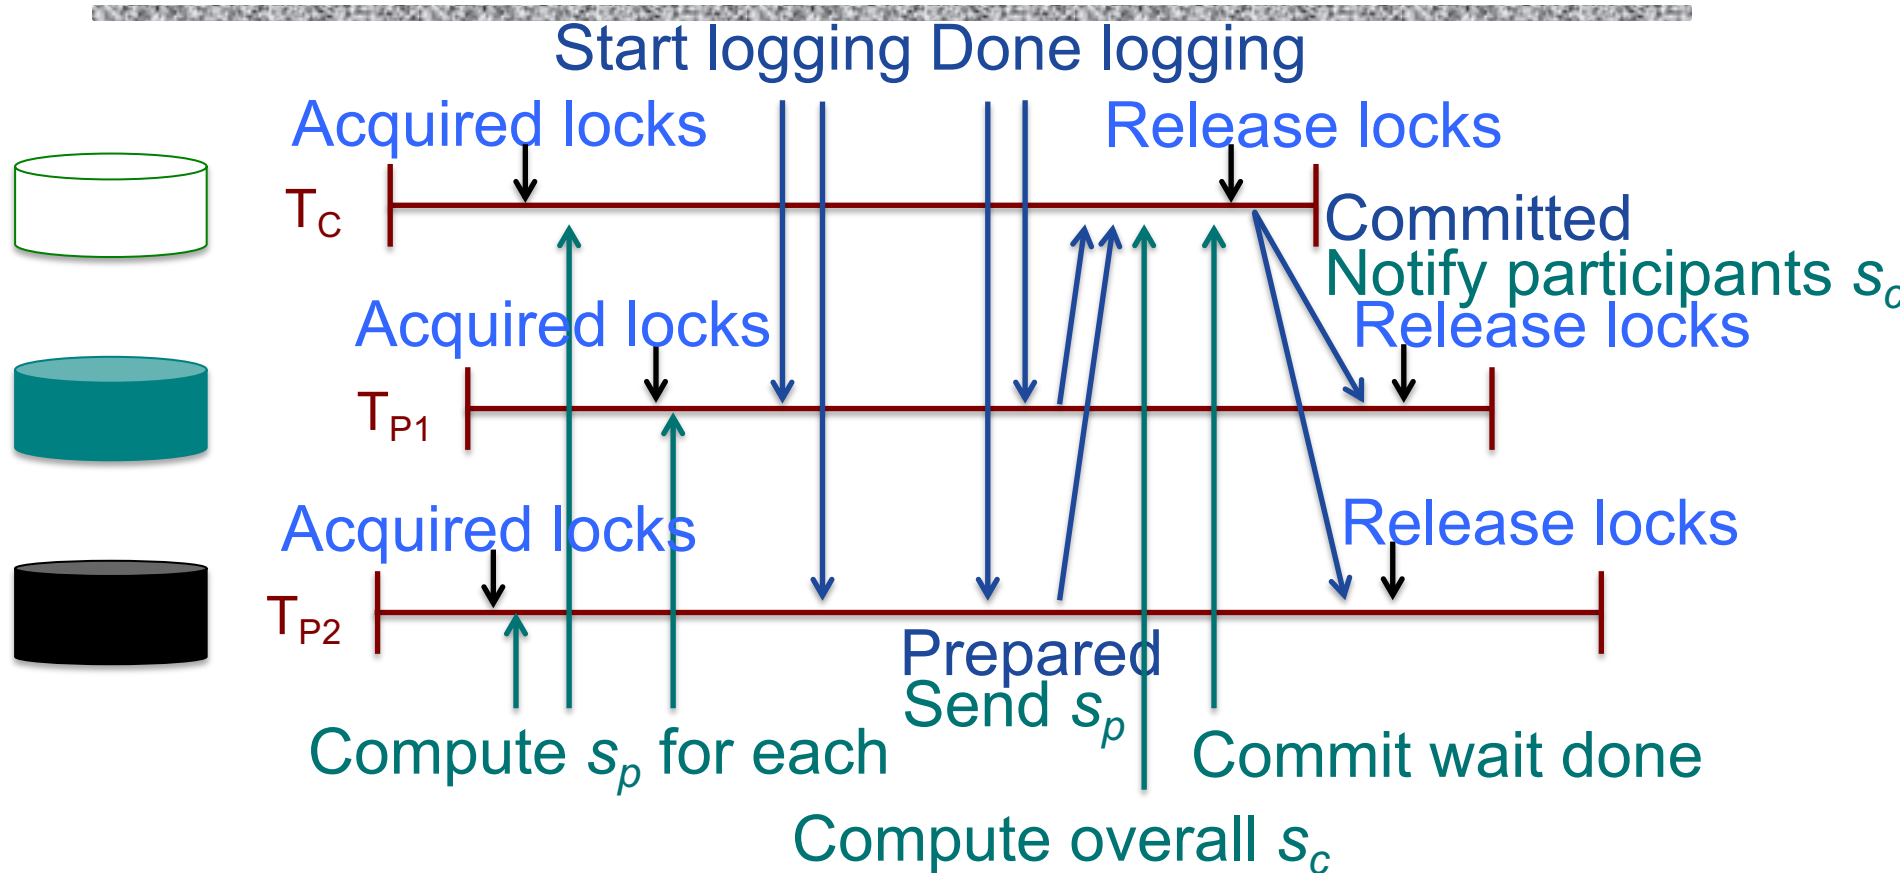
1. Client issues reads to leader of each tablet group, leader acquires read locks and returns most recent data

Commit Wait and 2-Phase Commit



2. Client locally performs (buffers) writes
3. ... chooses coordinator from leaders, initiates commit
4. ... sends commit msg to each leader (+ coord. identity)

Commit Wait and 2-Phase Commit

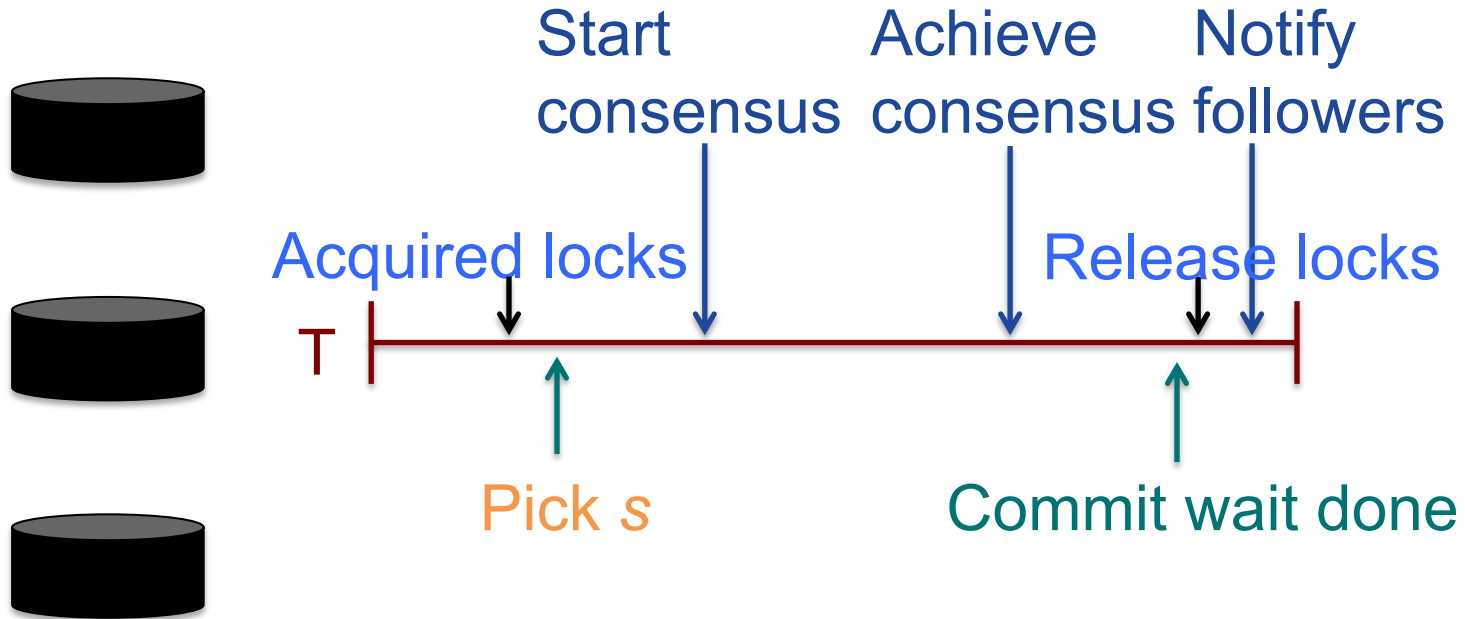


5. Client waits for commit from coordinator

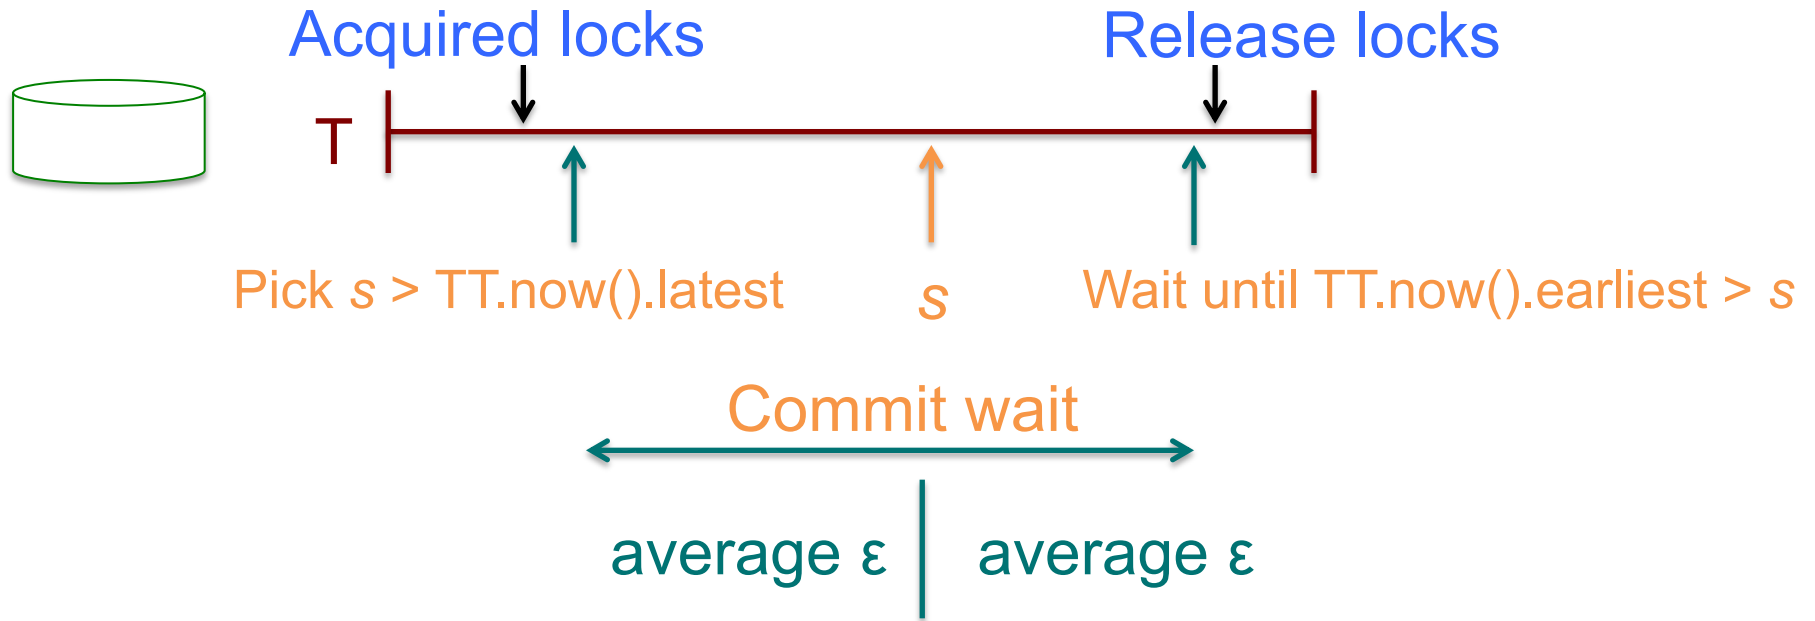
Commit Wait and 2-Phase Commit

- ◆ On commit msg from client, leaders acquire local write locks
 - If non-coordinator:
 - Choose prepare ts $>$ previous local timestamps
 - Log prepare record through Paxos
 - Notify coordinator of prepare timestamp
 - If coordinator:
 - Wait until hear from other participants
 - Choose commit timestamp \geq prepare ts $>$ local ts
 - Logs commit record through Paxos
 - Wait for the commit-wait period
 - Sends commit timestamp to replicas, other leaders, client
- ◆ All apply at-commit timestamp and release locks


Commit Wait and Replication






Commit Wait and TrueTime



Example



	Time	<8	8	15
 My friends		[X]	[]	
 My posts				[P]
 X's friends		[me]	[]	

Read-only Transactions (1)

- ◆ Assigning a timestamp requires a negotiation between all Paxos groups involved in the reads
 - Spanner requires a scope expression for every read-only transaction, summarizes the keys that will be read by the entire transaction
- ◆ If the scope's values are served by a single Paxos group, then the client issues the read-only transaction to that group's leader, who assigns reads and executes the transaction

Read-only Transactions (2)

If the scope's values are served by multiple Paxos groups...

- ◆ Option 1: Do a round of communication with all leaders to negotiate s_{read} based on `LastTS()`
- ◆ Option 2: The client avoids a negotiation round and has its reads execute at $s_{\text{read}} = \text{TT.now().latest}$
 - May wait for safe time to advance
 - All reads in the transaction can be sent to replicas that are sufficiently up-to-date

Schema-Change Transactions

- ◆ Schema-change transactions explicitly assigned a timestamp in the future
 - Future timestamp registered in the prepare phase
 - Atomic schema changes across thousands of servers do not disrupt concurrent activities
- ◆ Reads and writes, which implicitly depend on the schema, synchronize with any registered schema-change timestamp at time t
 - May proceed if their timestamps precede t
 - Must block behind the schema-change transaction if their timestamps are after t

Consistency Properties

◆ Eventual consistency

- Client may see an arbitrary subset of updates

◆ Snapshot isolation

- Client can read a consistent database snapshot at any time, incl. all updates up to snapshot timestamp
 - ... from any replica, in Spanner
 - Newer updates may still be replicating

◆ Strong consistency

- Client always sees the latest, consistent updates

Partition Tolerance in Spanner

- ◆ For reads?
- ◆ For writes?
- ◆ The role of TrueTime?

Remember the CAP theorem?