


Transport Layer



Computer Networking: A Top Down Approach

7th edition

Jim Kurose, Keith Ross

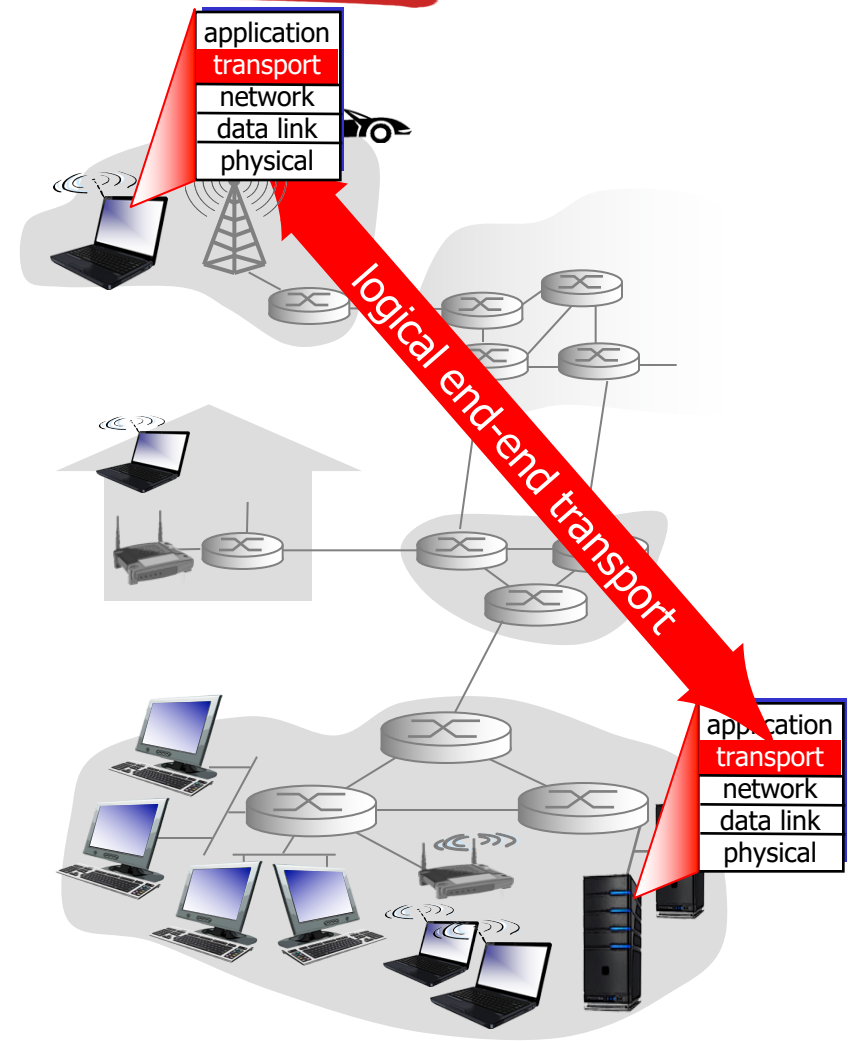
Pearson, 2017

Transport layer

- ❖ UDP: connectionless transport, not much to say
- ❖ TCP: connection-oriented reliable transport
- ❖ TCP congestion control

Transport services and protocols

- ❖ provide *logical communication* between app processes running on different hosts
- ❖ transport protocols run in end systems
 - send side: breaks app messages into *segments*, passes to network layer
 - rcv side: reassembles segments into messages, passes to app layer
- ❖ more than one transport protocol available to apps
 - Internet: TCP and UDP
- ❖ no delay/service guarantees



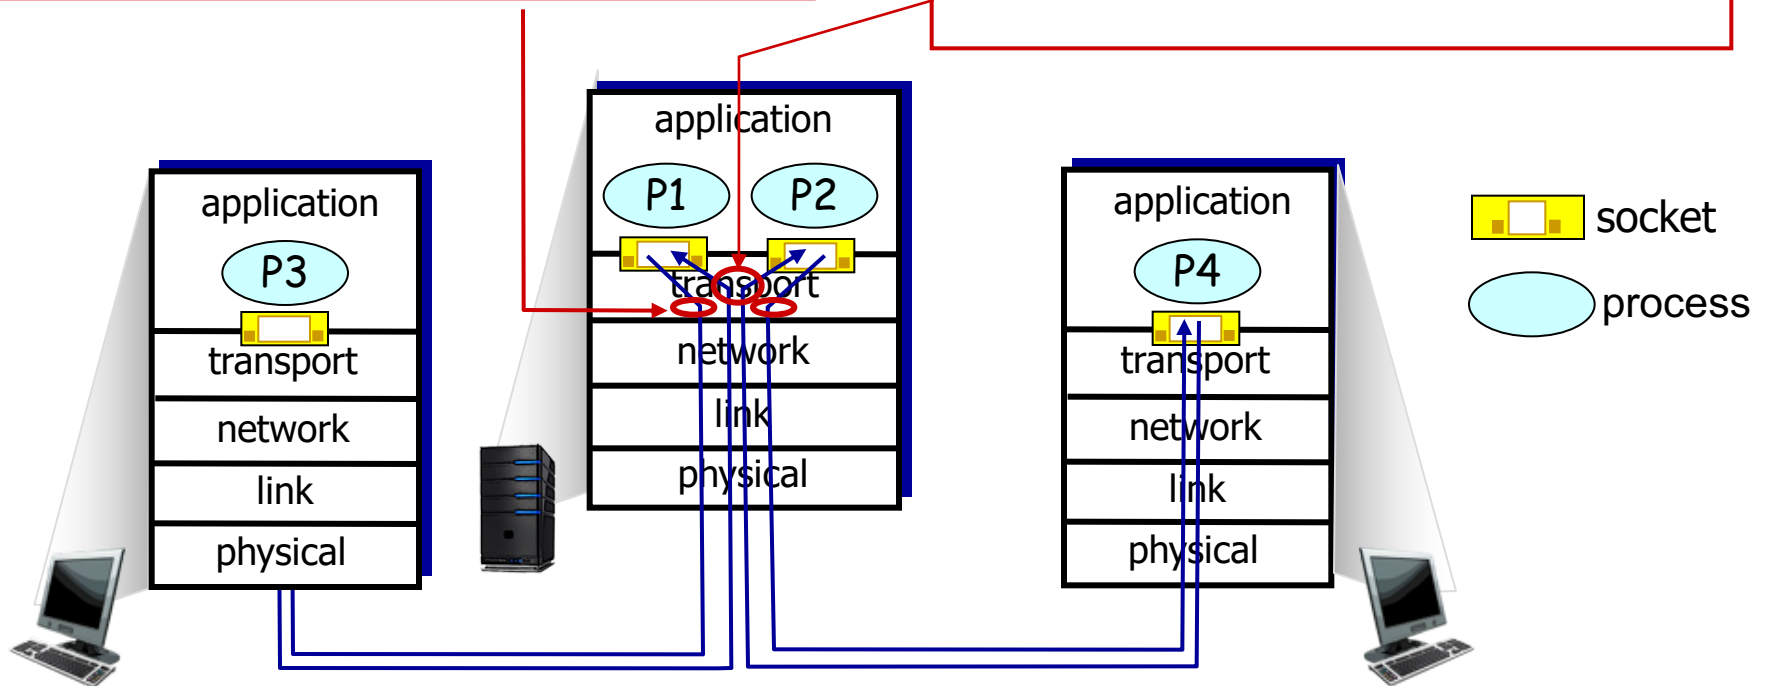
Multiplexing/demultiplexing

multiplexing at sender:

handle data from multiple sockets, add transport header (later used for demultiplexing)

demultiplexing at receiver:

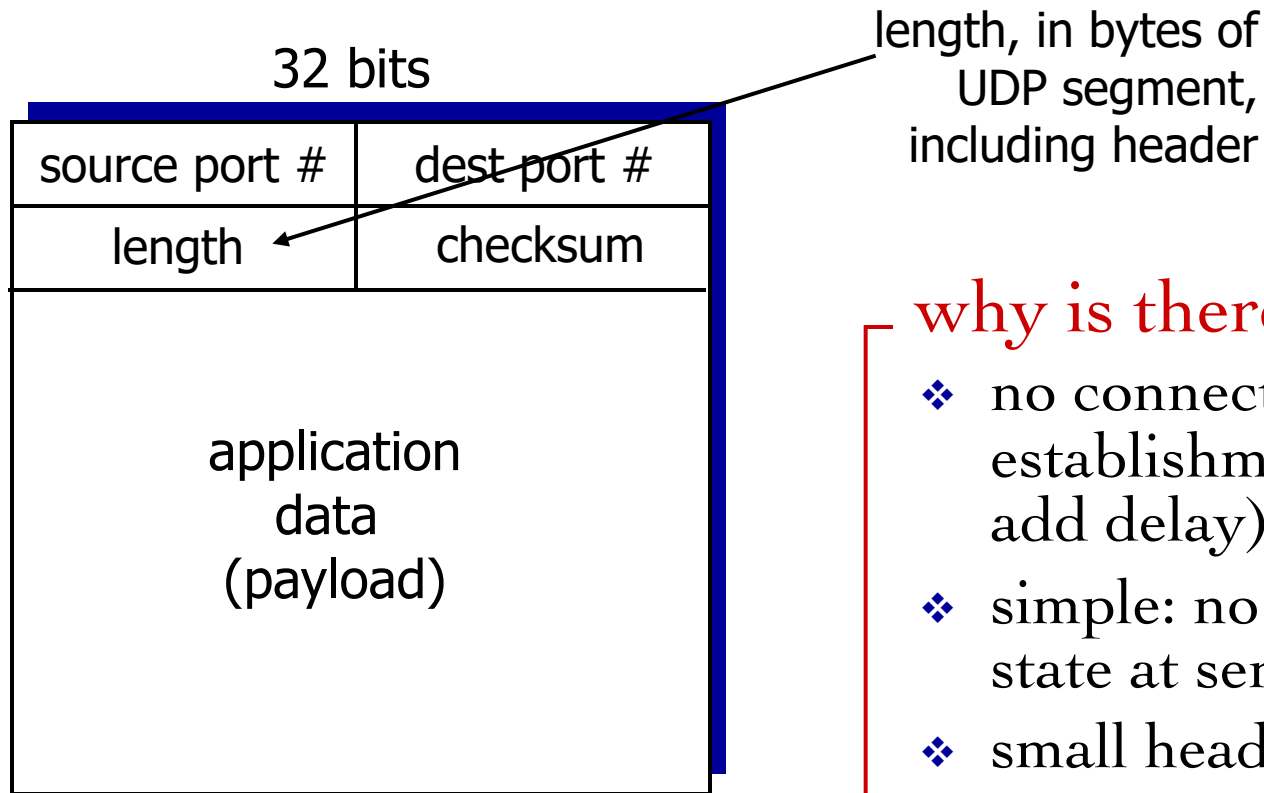
use header info to deliver received segments to correct socket



UDP: User Datagram Protocol [RFC 768]

- ❖ “no frills,” “bare bones”
Internet transport protocol
- ❖ “best effort” service, UDP segments may be:
 - lost
 - delivered out-of-order to app
- ❖ *connectionless*:
 - no handshaking between UDP sender, receiver
 - each UDP segment handled independently of others
- ❖ UDP use:
 - streaming multimedia apps (loss tolerant, rate sensitive)
 - DNS
 - SNMP
- ❖ reliable transfer over UDP:
 - add reliability at application layer
 - application-specific error recovery!

UDP segment header



UDP segment format

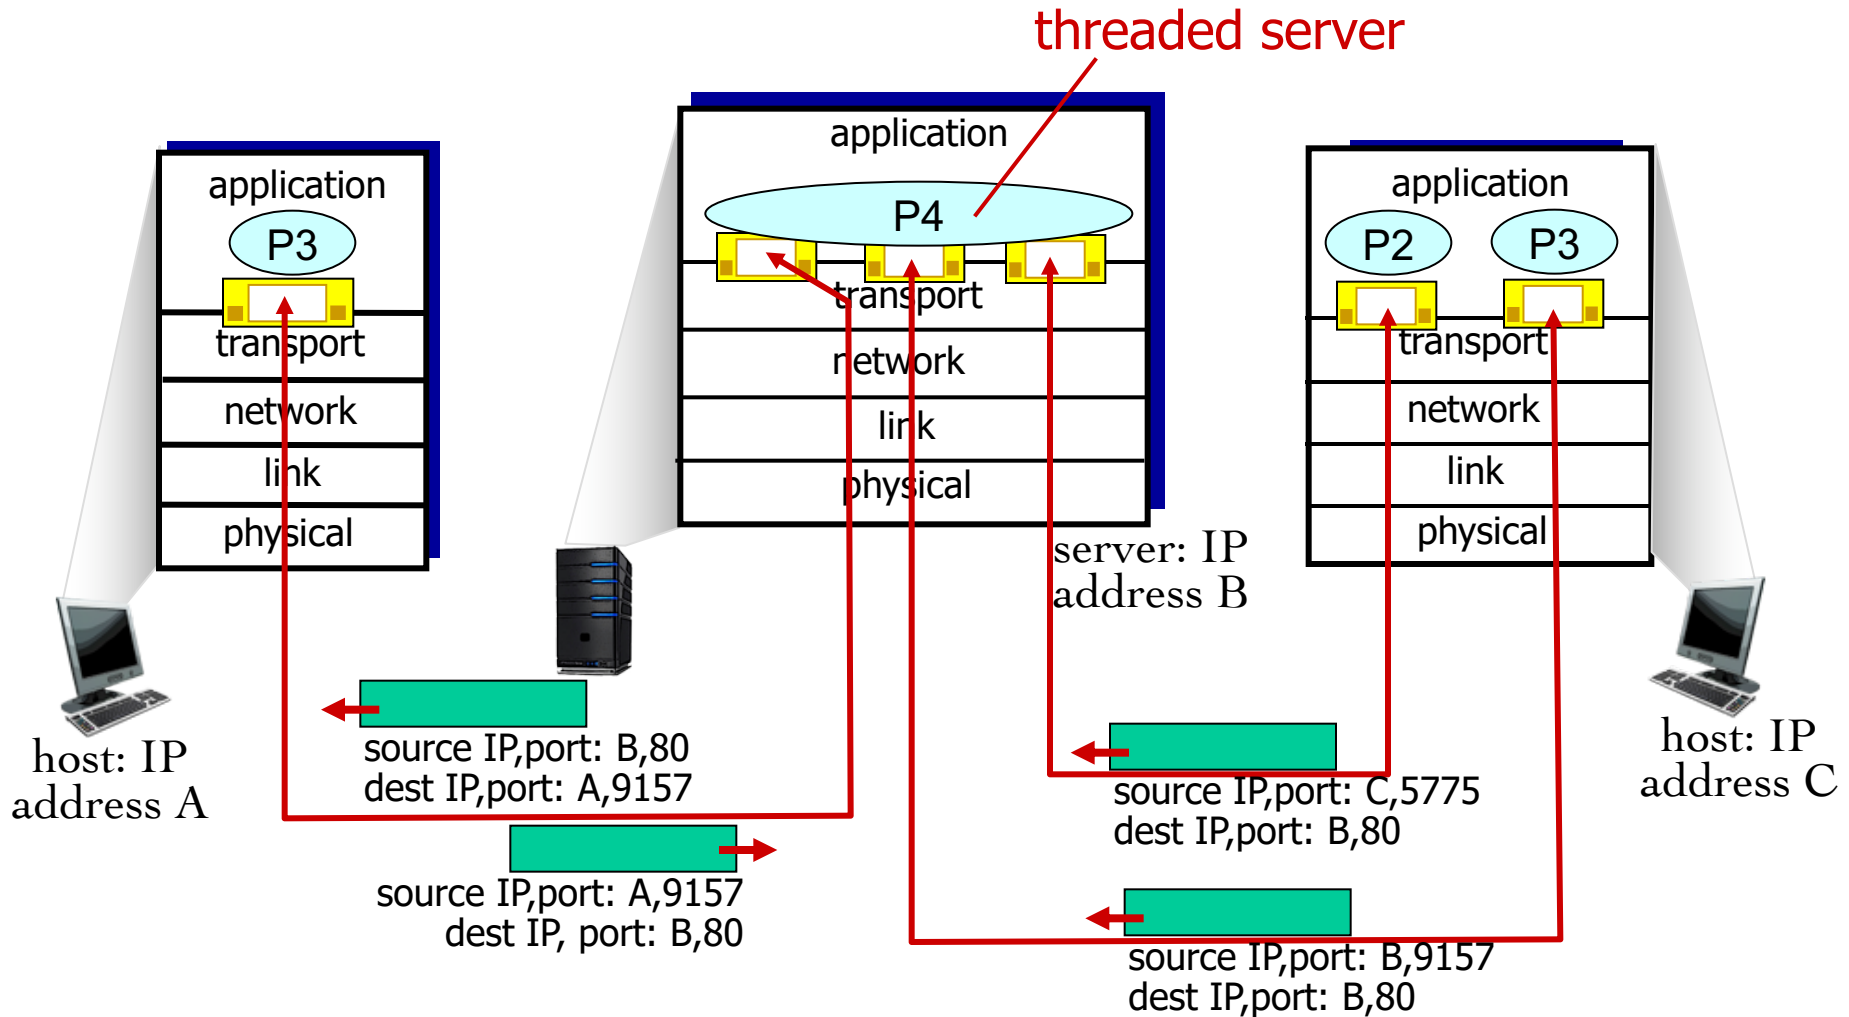
why is there a UDP?

- ❖ no connection establishment (which can add delay)
- ❖ simple: no connection state at sender, receiver
- ❖ small header size
- ❖ no congestion control: UDP can blast away as fast as desired

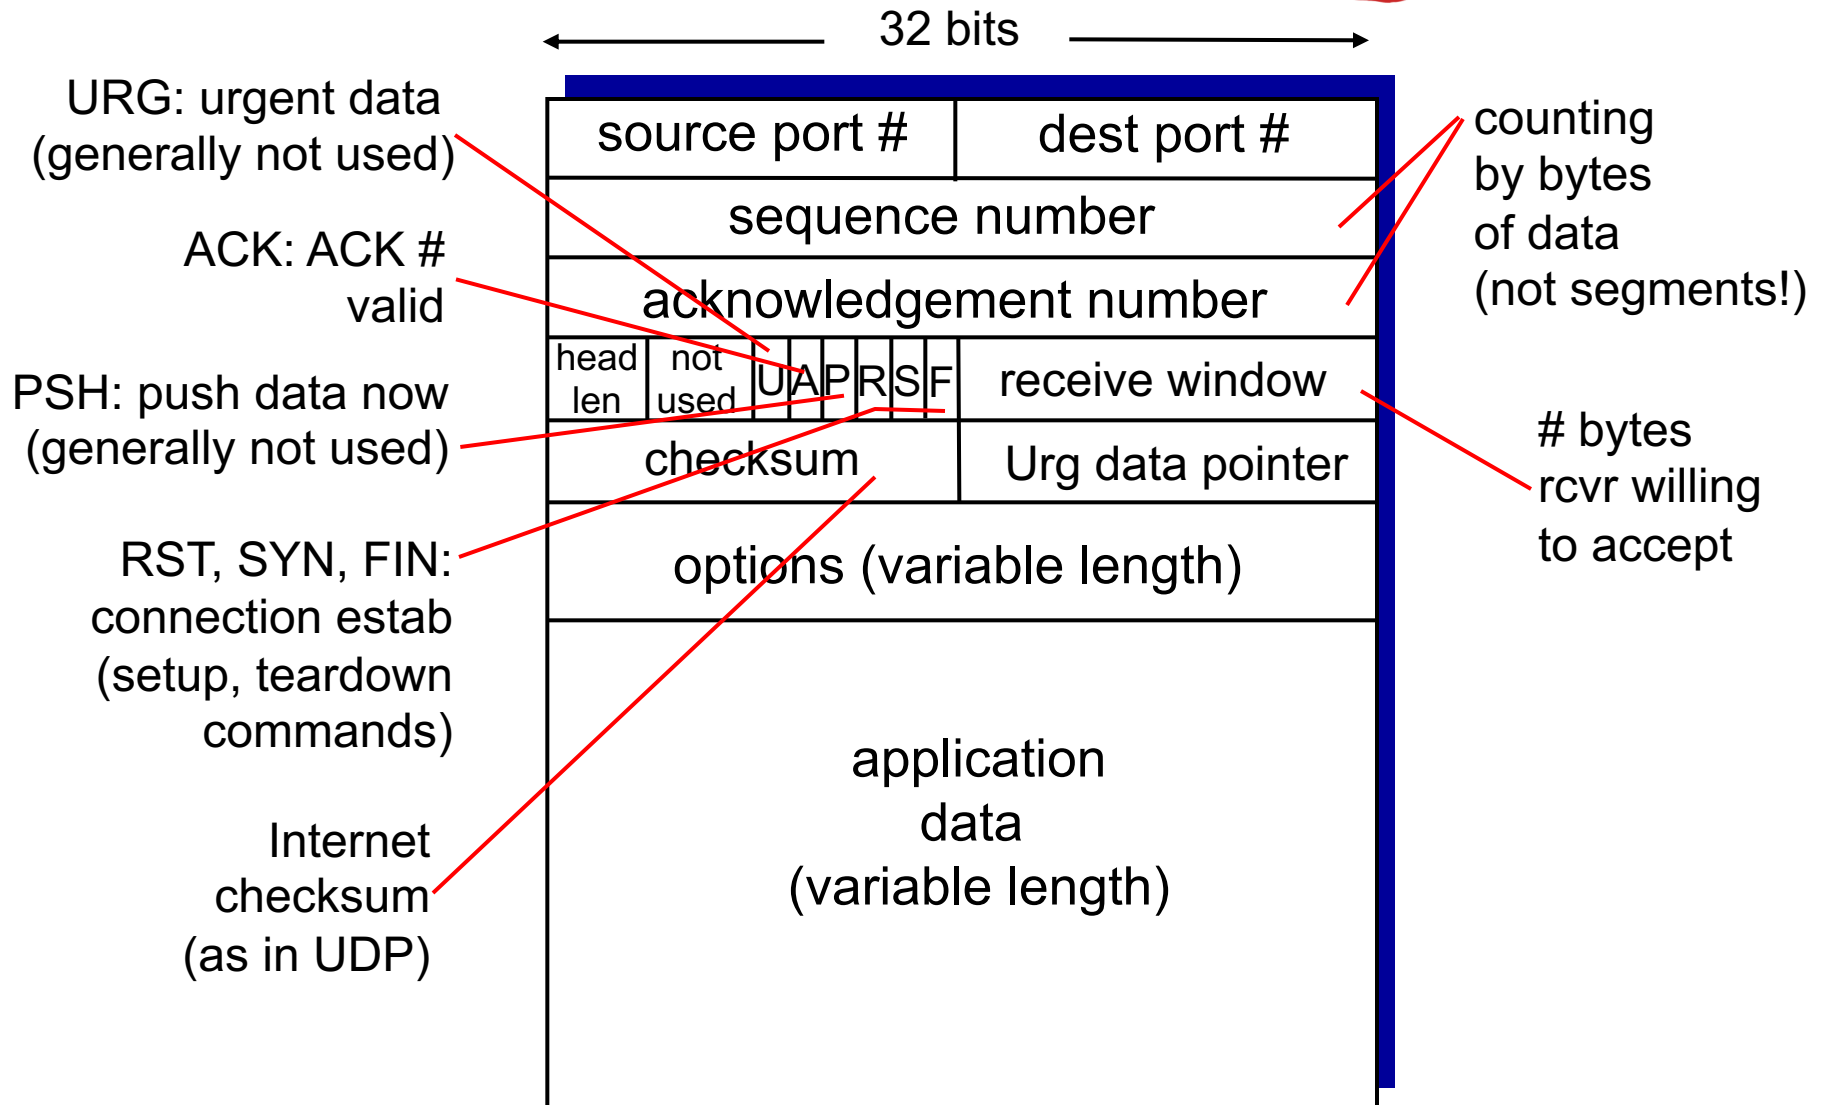
Connection-oriented demux

- ❖ TCP socket identified by 4-tuple:
 - source IP address
 - source port number
 - dest IP address
 - dest port number
- ❖ demux: receiver uses all four values to direct segment to appropriate socket
- ❖ server host may support many simultaneous TCP sockets:
 - each socket identified by its own 4-tuple
- ❖ web servers have different sockets for each connecting client
 - non-persistent HTTP will have different socket for each request

Connection-oriented: demux



TCP segment structure



TCP overview [RFCs 793,1122,1323, 2018, 2581]

❖ point-to-point:

- one sender, one receiver

❖ reliable, in-order *byte stream*:

- no “message boundaries”

❖ pipelined:

- TCP congestion and flow control set window size
- Hybrid of go back n and selective repeat

❖ full duplex data:

- bi-directional data flow in same connection
- MSS: maximum segment size

❖ connection-oriented:

- handshaking (exchange of control msgs) initializes sender, receiver state before data exchange

❖ flow controlled:

- sender will not overwhelm receiver

Channels with errors and loss

assumption: underlying channel can lose packets (data, ACKs)

- checksum, seq. #, ACKs, retransmissions will be of help ... but not enough


approach: sender waits “reasonable” amount of time for ACK

- ❖ retransmits if no ACK received in this time
- ❖ if pkt (or ACK) just delayed (not lost):
 - retransmission will be duplicate, but seq. #'s already handles this
 - receiver must specify seq # of pkt being ACKed
- ❖ requires countdown timer

Pipelined protocols

pipelining: sender allows multiple, “in-flight,” yet-to-be-acknowledged packets

- range of sequence numbers must be increased
- buffering at sender and/or receiver



(a) a stop-and-wait protocol in operation

(b) a pipelined protocol in operation

❖ two generic forms of pipelined protocols:
go-Back-N, selective repeat

Pipelined protocols: concepts

Go-back-N:

- ❖ sender can have up to N unack'ed packets in pipeline
- ❖ receiver only sends *cumulative ack*
 - Doesn't ack packet if there's a gap
- ❖ sender has timer for oldest unacked packet
 - when timer expires, retransmit *all* unacked packets

Selective Repeat:

- ❖ sender can have up to N unack'ed packets in pipeline
- ❖ rcvr sends *individual ack* for each packet
- ❖ sender maintains timer for each unacked packet
 - when timer expires, retransmit only that unacked packet

TCP sequence numbers and ACKs

sequence numbers:

- byte stream “number” of first byte in segment’s data

acknowledgements:


- seq # of next byte expected from other side
- cumulative ACK

Q: how receiver handles out-of-order segments

- A:** TCP spec doesn’t say,
- up to implementor

outgoing segment from sender

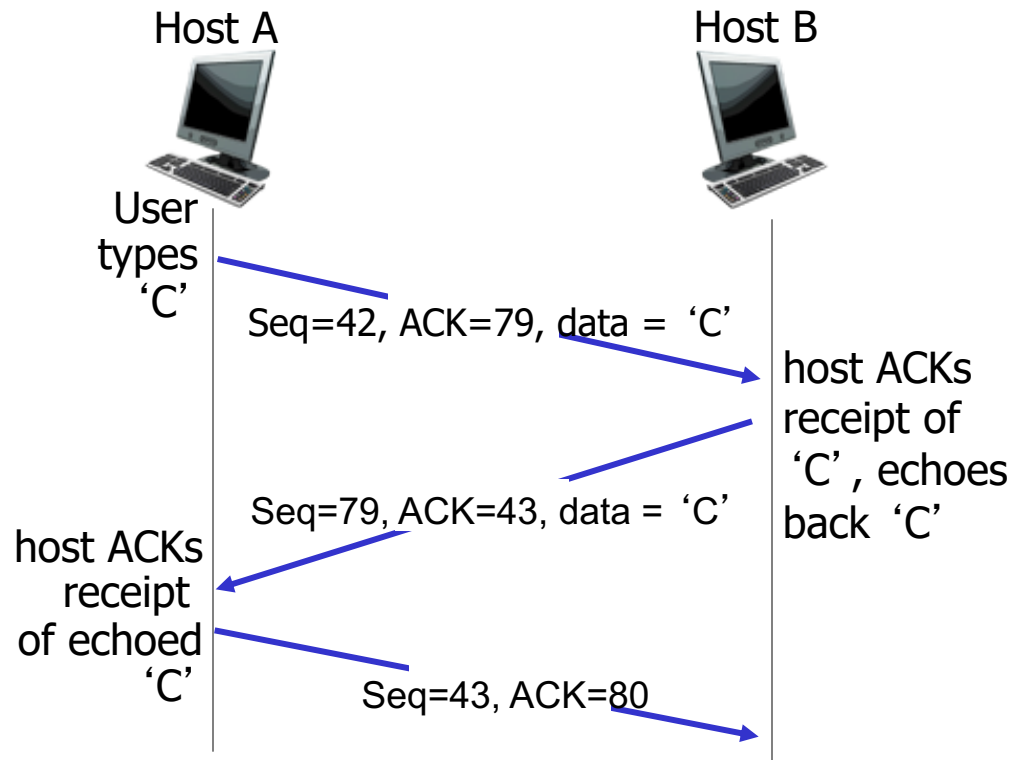
source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer



incoming segment to sender

source port #	dest port #
sequence number	
acknowledgement number	
	A
	rwnd
checksum	urg pointer

TCP sequence numbers and acks



simple telnet scenario

TCP RTT estimation

Q: how to set TCP timeout value?

- ❖ longer than RTT
 - but RTT varies
- ❖ *too short*: premature timeout, unnecessary retransmissions
- ❖ *too long*: slow reaction to segment loss


Q: how to estimate RTT?

- ❖ **SampleRTT**: measured time from segment transmission until ACK receipt
 - ignore retransmissions
- ❖ **SampleRTT** will vary, want estimated RTT “smoother”
 - average several recent measurements, not just current **SampleRTT**

TCP RTT estimation

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- ❖ exponential weighted moving average
- ❖ influence of past sample decreases exponentially fast
- ❖ typical value: $\alpha = 0.125$



TCP timeout

- ❖ **timeout interval**: **EstimatedRTT** plus “safety margin”
 - large variation in **EstimatedRTT** -> larger safety margin
- ❖ estimate SampleRTT deviation from EstimatedRTT:

$$\text{DevRTT} = (1-\beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically, $\beta = 0.25$)

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$



↑
estimated RTT

↑
“safety margin”

TCP sender (no optimization, congestion control)

data rcvd from app:

- ❖ create segment with seq #
- ❖ seq # is byte-stream number of first data byte in segment
- ❖ start timer if not already running
 - think of timer as for oldest unacked segment
 - expiration interval: **TimeoutInterval**

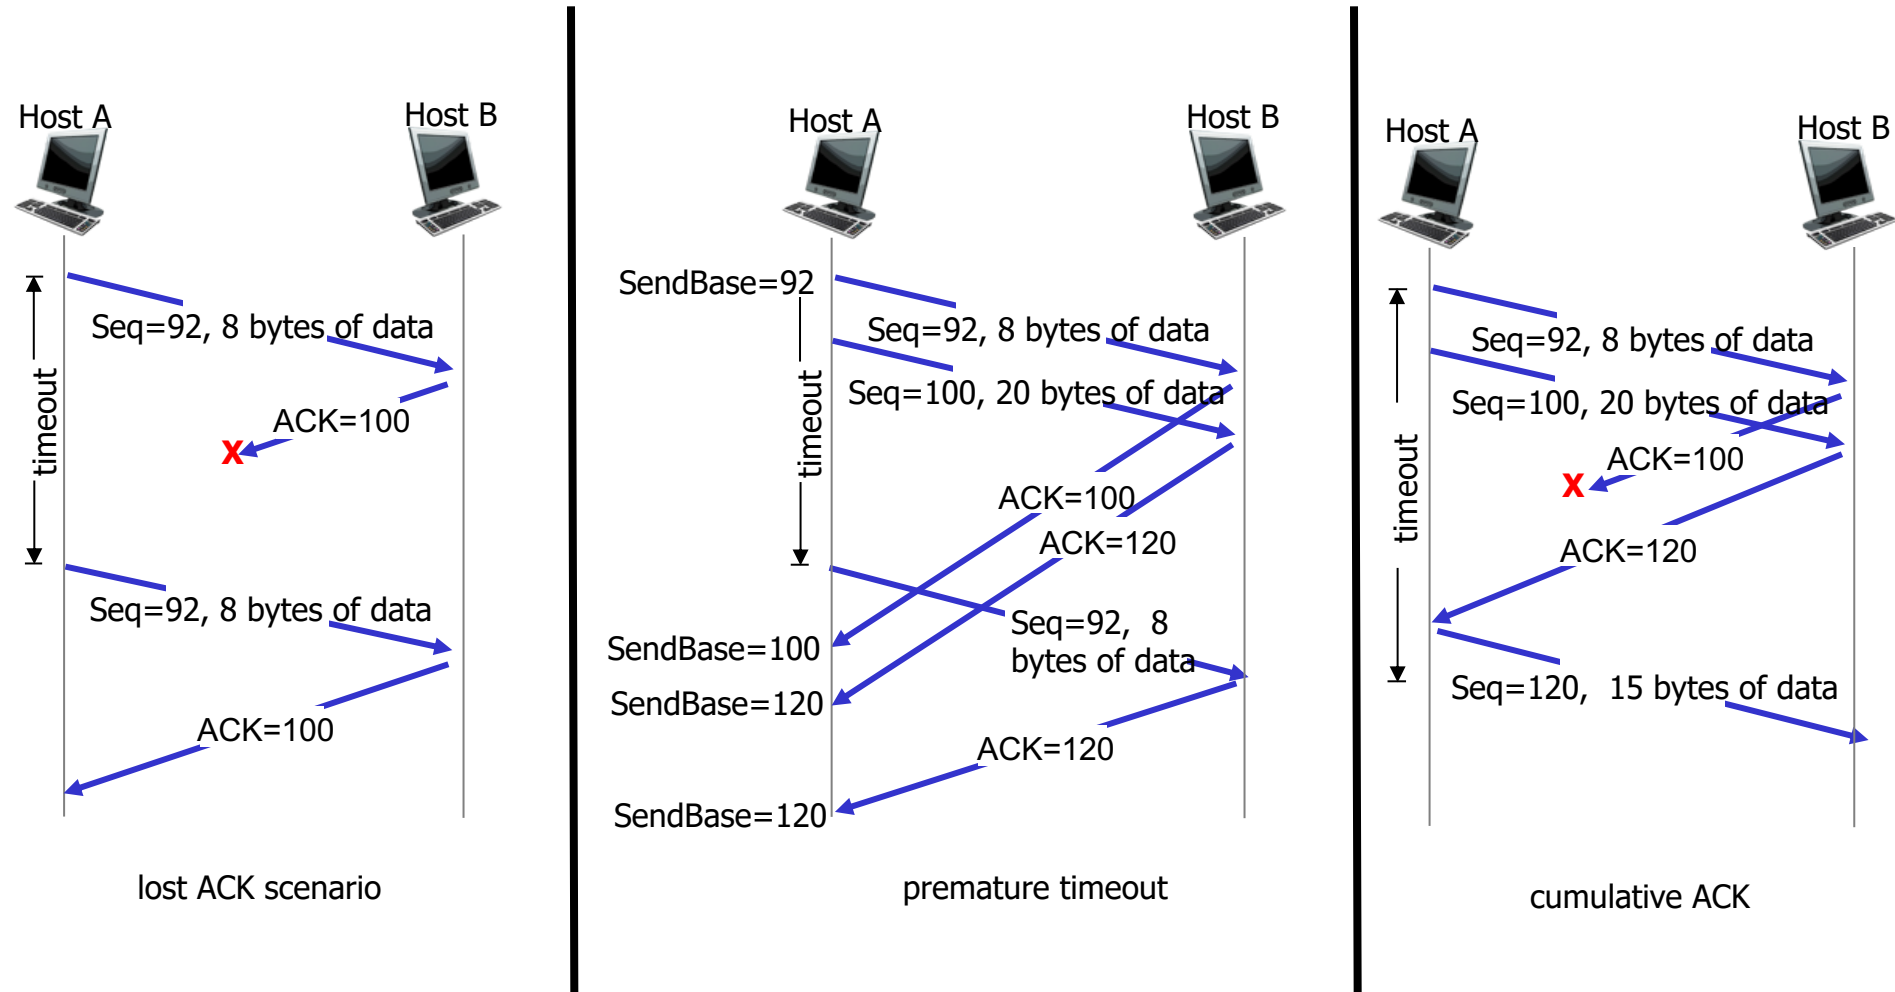
timeout:

- ❖ retransmit segment that caused timeout
- ❖ restart timer

ack rcvd:

- ❖ if ack acknowledges previously unacked segments
 - update what is known to be ACKed
 - start timer if there are still unacked segments

TCP retransmission scenarios



TCP ACK original generation [RFC 1122, 2581]

<i>event at receiver</i>	<i>TCP receiver action</i>
arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	delayed ACK: wait up to 500ms for next segment, if no next segment, send ACK
arrival of in-order segment with expected seq #. One other segment has ACK pending	immediately send single cumulative ACK, ACKing both in-order segments
arrival of out-of-order segment higher-than-expected seq. # . Gap detected	immediately send duplicate ACK , indicating seq. # of next expected byte
arrival of segment that partially or completely fills gap	immediate send ACK, provided that segment starts at lower end of gap

TCP fast retransmit

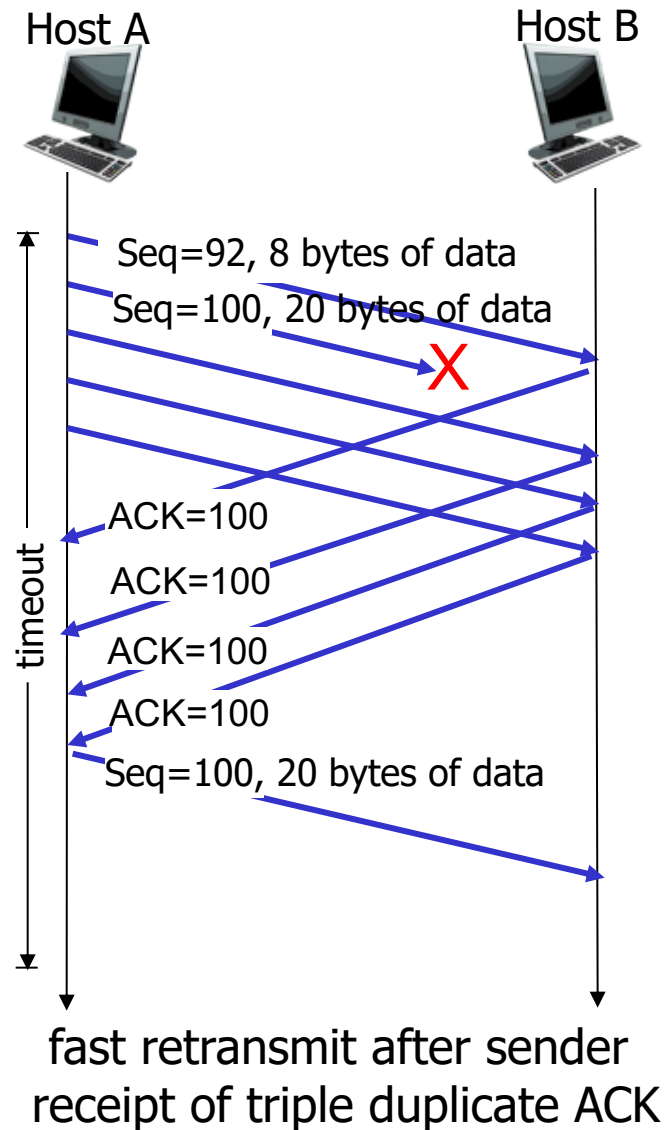
- ❖ time-out period often relatively long:
 - long delay before resending lost packet
- ❖ detect lost segments via duplicate ACKs.
 - sender often sends many segments back-to-back
 - if segment is lost, there will likely be many duplicate ACKs.

TCP fast retransmit

if sender receives 3 ACKs for same data (“triple duplicate ACKs”), resend unacked segment with smallest seq #

- likely that unacked segment lost, so don't wait for timeout


TCP fast retransmit



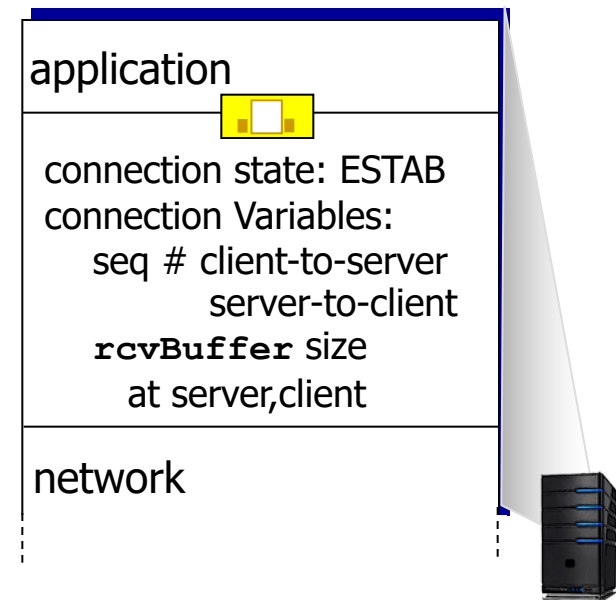
Connection management

before exchanging data, sender/receiver “handshake”:

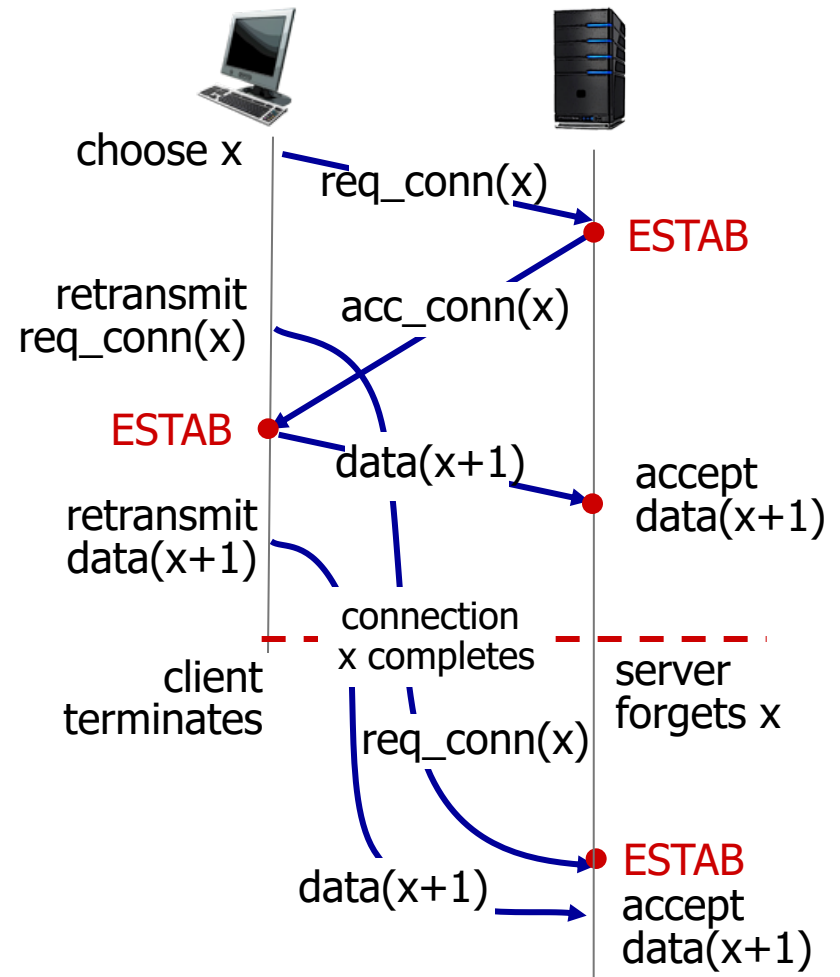
- ❖ agree to establish connection (each knowing the other willing to establish connection)
- ❖ agree on connection parameters – sequence number



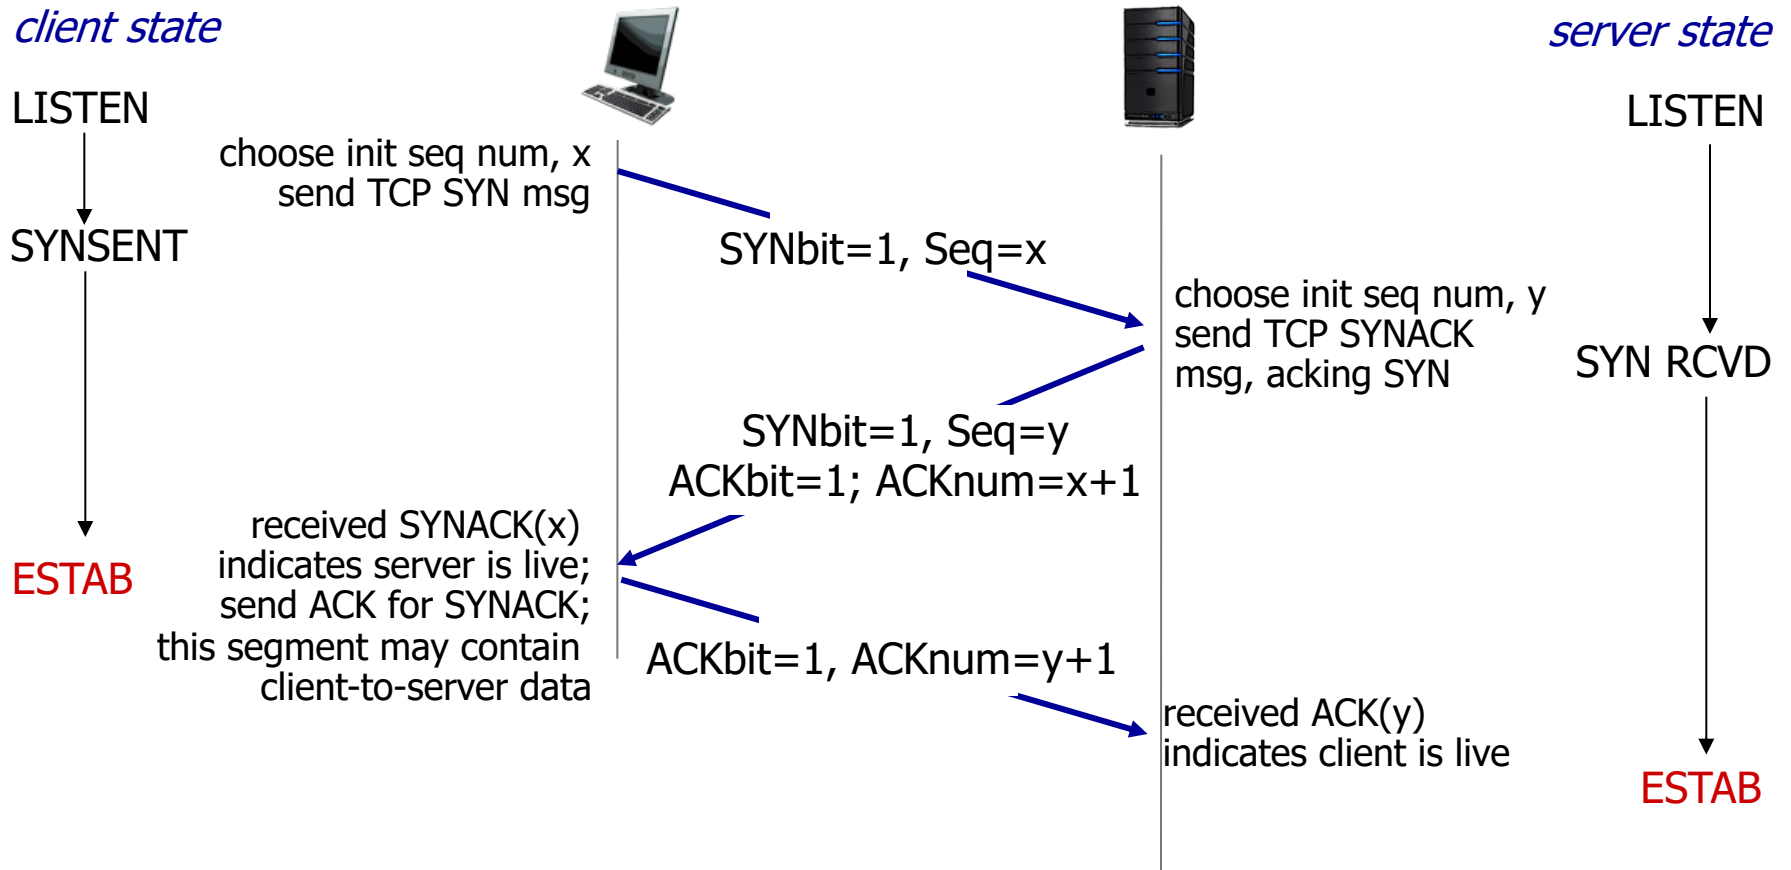
```
Socket clientSocket =  
    newSocket("hostname", "port number");
```



```
Socket connectionSocket =  
    welcomeSocket.accept();
```



TCP 3-way handshake



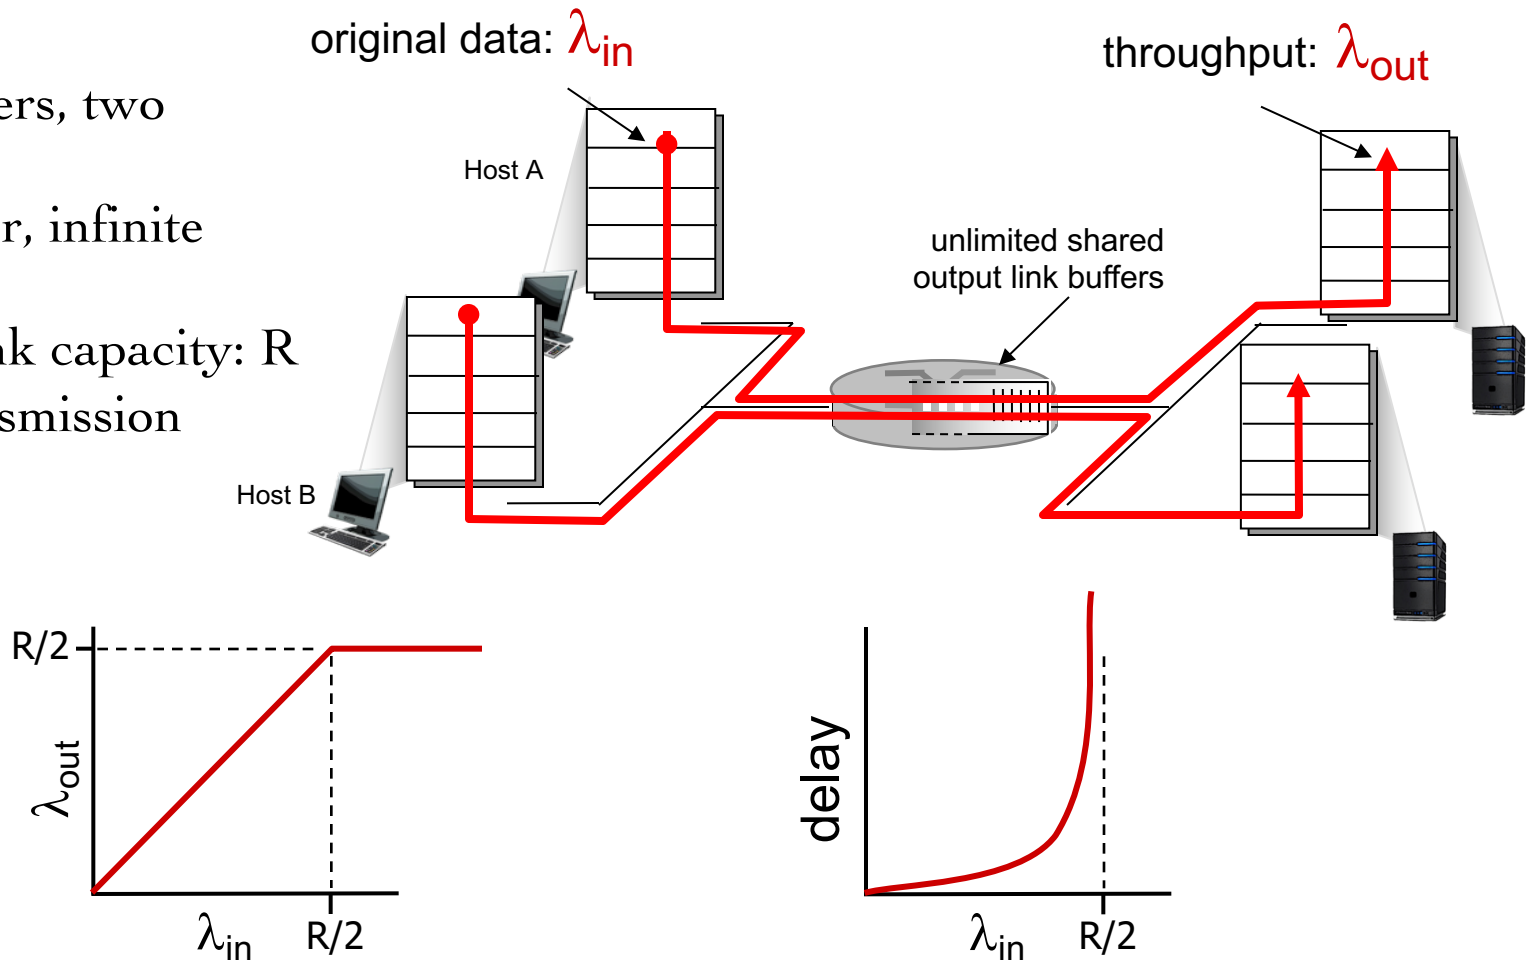
Principles of congestion control

congestion:

- ❖ informally: “too many sources sending too much data too fast for *network* to handle”
- ❖ different from flow control!
- ❖ manifestations:
 - lost packets (buffer overflow at routers)
 - long delays (queueing in router buffers)
- ❖ a top-10 problem!

Causes/costs of congestion: delay


- ❖ two senders, two receivers
- ❖ one router, infinite buffers
- ❖ output link capacity: R
- ❖ no retransmission



- ❖ maximum per-connection throughput: $R/2$
- ❖ large delays as arrival rate, λ_{in} , approaches capacity

Causes/costs of congestion: drop + duplicate


- ❖ one router, *finite* buffers
- ❖ sender retransmission of timed-out packet
 - application-layer input = application-layer output: $\lambda_{in} = \lambda_{out}$
 - transport-layer input includes *retransmissions*: $\lambda'_{in} \geq \lambda_{in}$



Causes/costs of congestion

duplicates



- ❖ packets can be lost, dropped at router due to full buffers
- ❖ sender times out prematurely, sending **two** copies, both of which are delivered



costs of congestion:

- ❖ more work (retransmissions) for given “goodput”
- ❖ unneeded retransmissions: link carries multiple copies of pkts
 - decreasing goodput

Causes/costs of congestion: multi-hop paths



another “cost” of congestion:

- ❖ when packet dropped, any “upstream” transmission capacity used for that packet was wasted!

Approaches to congestion control

two broad approaches towards congestion control:

end-end congestion control:

- ❖ no explicit feedback from network
- ❖ congestion inferred from end-system observed loss, delay, Ack signals from receiver
- ❖ approach taken by TCP


network-assisted congestion control:

- ❖ routers provide feedback to end systems
 - single bit indicating congestion as packet forwarded (SNA, DECbit, TCP/IP ECN)


TCP congestion control basics

- ❖ **approach**: sender increases transmission rate (window size), probing for usable bandwidth, until loss occurs
 - **additive increase**: increase **cwnd** by 1 MSS (max segment size, default 536 bytes) every RTT until loss detected
 - **multiplicative decrease**: cut **cwnd** in half after loss

AIMD saw tooth behavior: probing for bandwidth



TCP congestion control window



- ❖ sender limits transmission:

$$\text{LastByteSent} - \text{LastByteAcked} \leq \text{cwnd}$$

- ❖ **cwnd** is dynamic, function of perceived network congestion


TCP sending rate:

- ❖ roughly... send cwnd bytes, wait RTT for ACKS, then send more bytes

$$\text{rate} \approx \frac{\text{cwnd}}{\text{RTT}} \text{ bytes/sec}$$

TCP slow start

- ❖ when connection begins, rate (window) small but increase rate (window) exponentially until first loss event:
 - initially **cwnd** = 1 MSS
 - double **cwnd** every RTT
 - done by incrementing **cwnd** for every ACK received
- ❖ summary: initial rate is slow but ramps up exponentially fast



TCP: detecting loss

- ❖ Timeout
- ❖ Duplicate ACKs from receiver
 - Receiver adjusted to **send ACK each time packet received** even if it's a duplicate (no advance in seq number)
 - Sender waits for **3 duplicate ACKs** (not just one in case slightly out of order delivery)

TCP: responding to loss

- ❖ If loss detected by timeout, **cwnd** set to 1 MSS
 - window initially grows exponentially to threshold (ssthresh), then grows linearly
- ❖ If lost detected by duplicate ACKS, can be less reactive since ACKs indicate network able to deliver packets to receiver
 - **TCP Reno:** **cwnd** is only cut in half (not set to 1); then window grows linearly
 - **TCP New Reno:** improves retransmit during fast recovery given wireless link loss is not usually due to network congestion
 - for every ACK that advances sequence space, send next packet beyond the ACKed sequence number as well


TCP: switching from slow start to CA

Q: when should the exponential increase switch to linear?

A: when **cwnd** gets to 1/2 of its value before timeout.

Implementation:


- ❖ variable **ssthresh**
- ❖ on loss event, **ssthresh** is set to 1/2 of **cwnd** just before loss event




TCP throughput

- ❖ avg. TCP thruput as function of window size, RTT?
 - ignore transient of slow start, assume always data to send
- ❖ W : window size (measured in bytes) where loss occurs
 - avg. window size (# in-flight bytes) is $\frac{3}{4} W$
 - avg. thruput is $\frac{3}{4}W$ per RTT

$$\text{avg TCP thruput} = \frac{3}{4} \frac{W}{\text{RTT}} \text{ bytes/sec}$$




Understanding network behavior is hard



TCP fairness


fairness goal: if K TCP sessions share same bottleneck link of bandwidth R , each should have average rate of R/K



Why is TCP fair, in concept

two competing sessions:

- ❖ additive increase gives slope of 1, as throughput increases
- ❖ multiplicative decrease decreases throughput proportionally



Side note: proportional fair scheduling in 3G

```
[~, b_user] = max(drc(i, :)/_avg_thruput(i, :));  
avg_thruput(i+1, :) = (i/(i+1))*avg_thruput(i, :);  
avg_thruput(i+1, b_user) = (i/(i+1))*avg_thruput(i, b_user) +  
drc(i, b_user)/(i+1);
```

“Transmitter directed, multiple receiver system using path diversity to equitably maximize throughput,” U.S. Patent No. 6449490, Sept. 10, 2002